

Scale Invariant Feature Transform

Jenn-Jier James Lien (連 震 杰)
Professor

Computer Science and Information Engineering
National Cheng Kung University

(O) (06) 2757575 ext. 62540

jjlien@csie.ncku.edu.tw

<http://robotics.csie.ncku.edu.tw>

Major Issues

1. SIFT: Scale Invariant Feature Transform
2. Detection of Scale-Space Extrema
 - 1) Difference of Gaussian (DoG) filter Vs. Laplacian of Gaussian (LoG)
 - 2) Non-maxima/minima suppression (Local Maximum)
3. Keypoint Localization (or Feature Selection)
 - 1) Keypoint localization in sub-pixel accuracy
 - 2) Remove low contrast keypoints
 - 3) Eliminating edge responses by Hessian matrix
4. Keypoint/Feature Orientation Assignment
5. Keypoint/Feature Descriptor

Content

1. Motivation
2. System Flowchart
3. Detection of Scale-Space Extrema
 - 3.1 Difference of Gaussian (DoG) filter
 - 3.2 Non-maxima/minima suppression
4. Keypoint Localization
 - 4.1 Keypoint localization in sub-pixel accuracy
 - 4.2 Remove low contrast keypoints
 - 4.3 Eliminating edge responses
5. Orientation Assignment
6. Keypoint Descriptor
7. Parameter Selection

Content

1. Motivation
2. System Flowchart
3. Detection of Scale-Space Extrema
 - 3.1 Difference of Gaussian (DoG) filter
 - 3.2 Non-maxima/minima suppression
4. Keypoint Localization
 - 4.1 Keypoint localization in sub-pixel accuracy
 - 4.2 Remove low contrast keypoints
 - 4.3 Eliminating edge responses
5. Orientation Assignment
6. Keypoint Descriptor
7. Parameter Selection

Motivation

- This paper presents a method for **extracting distinctive invariant features** from images.
- The features are **invariant** to image **scale** and **rotation**.
- The features are shown to provide robust matching across a substantial range of **affine distortion**, **change in 3D viewpoint**, addition of **noise**, and **change in illumination**.

Content

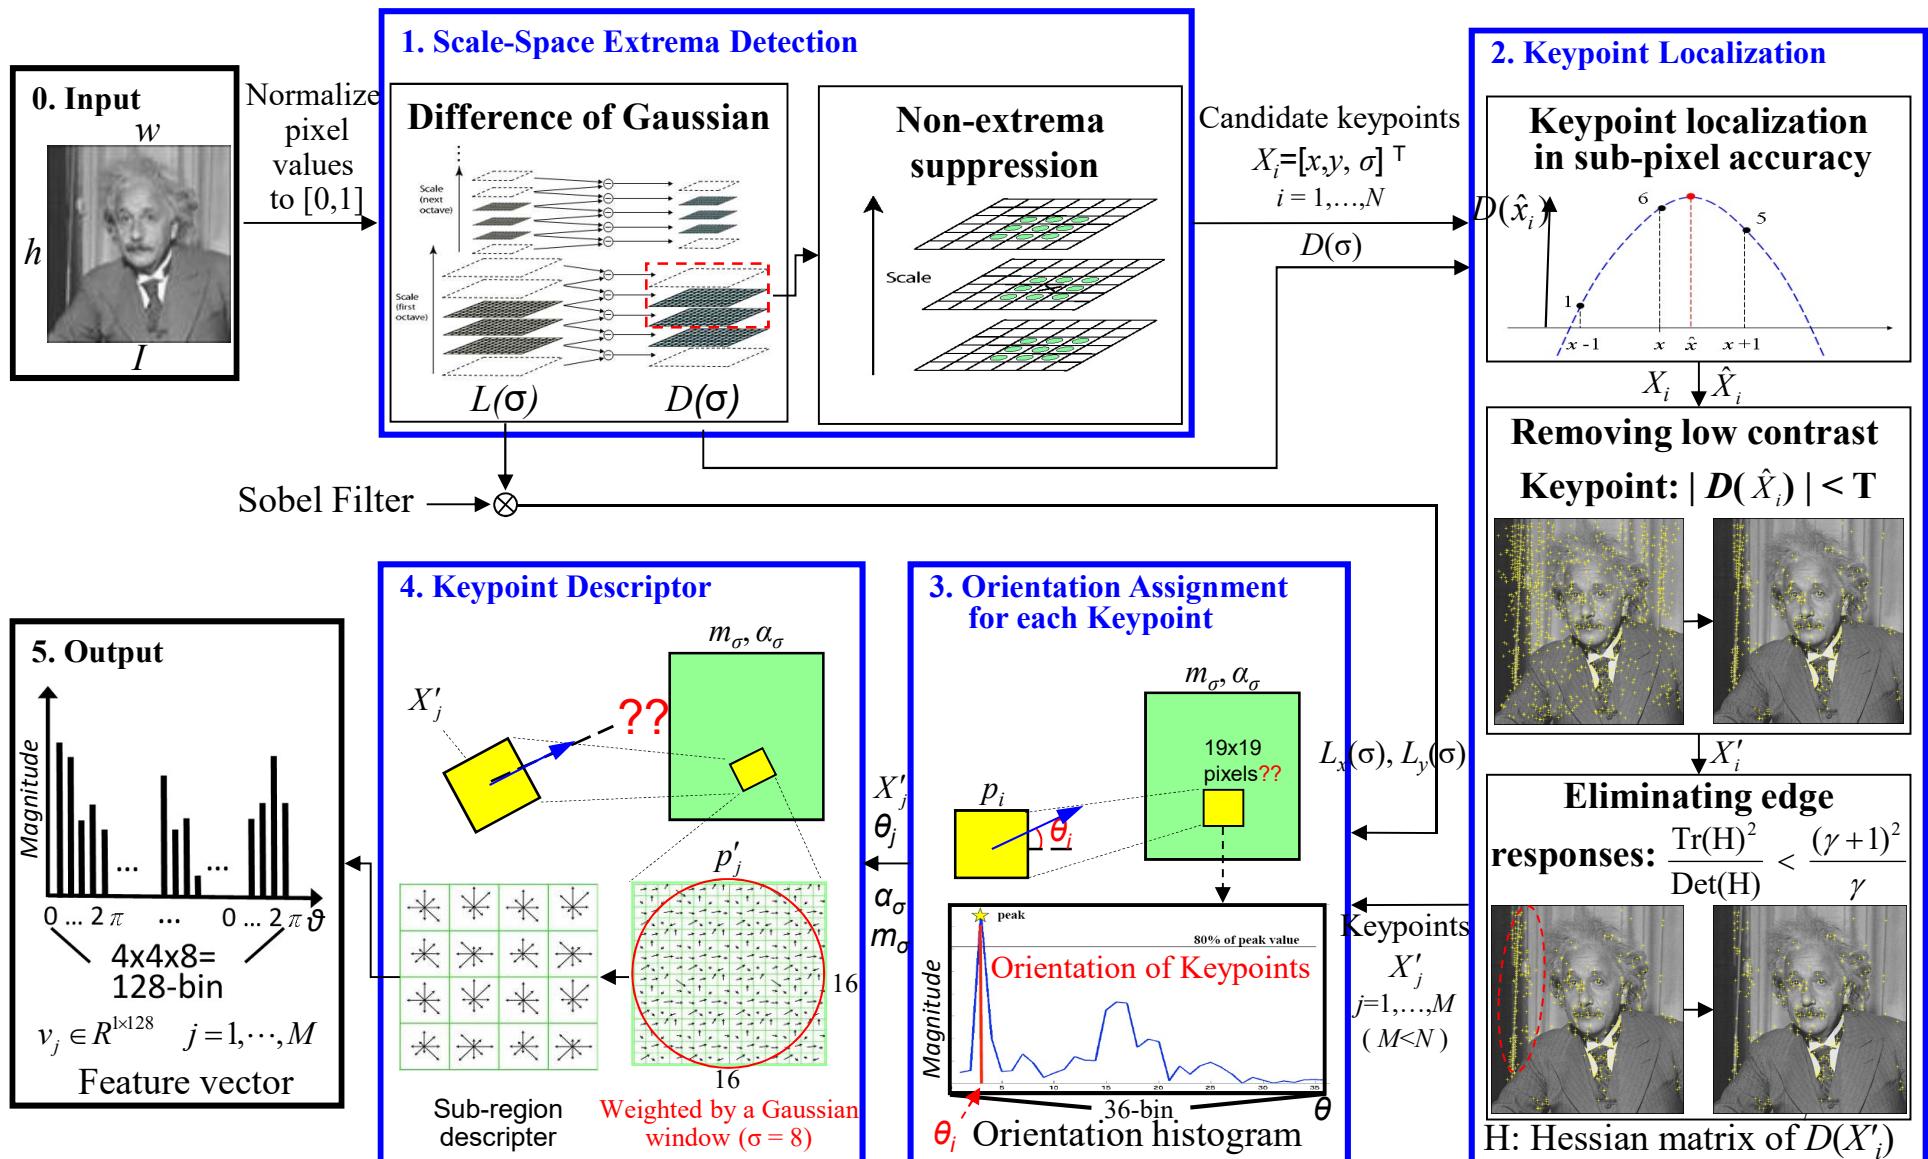
1. Motivation
2. System Flowchart
3. Detection of Scale-Space Extrema
 - 3.1 Difference of Gaussian (DoG) filter
 - 3.2 Non-maxima/minima suppression
4. Keypoint Localization
 - 4.1 Keypoint localization in sub-pixel accuracy
 - 4.2 Remove low contrast keypoints
 - 4.3 Eliminating edge responses
5. Orientation Assignment
6. Keypoint Descriptor
7. Parameter selection

$$L(\sigma) = G(\sigma) * I$$

$$D(\sigma) = (G(k^{(n-1)}\sigma) - G(k^n\sigma)) * I$$

(Scale space)	α_σ	(Gradient orientation of $L(\sigma)$)
(DoG images)	m_σ	(Gradient magnitude of $L(\sigma)$)
	X'_i	(Accurate keypoint localization)
	θ_j	(Orientation of i -th keypoint)

2. System Flowchart



Content

1. Motivation

2. System Flowchart

3. Detection of Scale-Space Extrema

 3.1 Difference of Gaussian (DoG) filter

 3.2 Non-maxima/minima suppression

4. Keypoint Localization

 4.1 Keypoint localization in sub-pixel accuracy

 4.2 Remove low contrast keypoints

 4.3 Eliminating edge responses

5. Orientation Assignment

6. Keypoint Descriptor

7. Parameter Selection

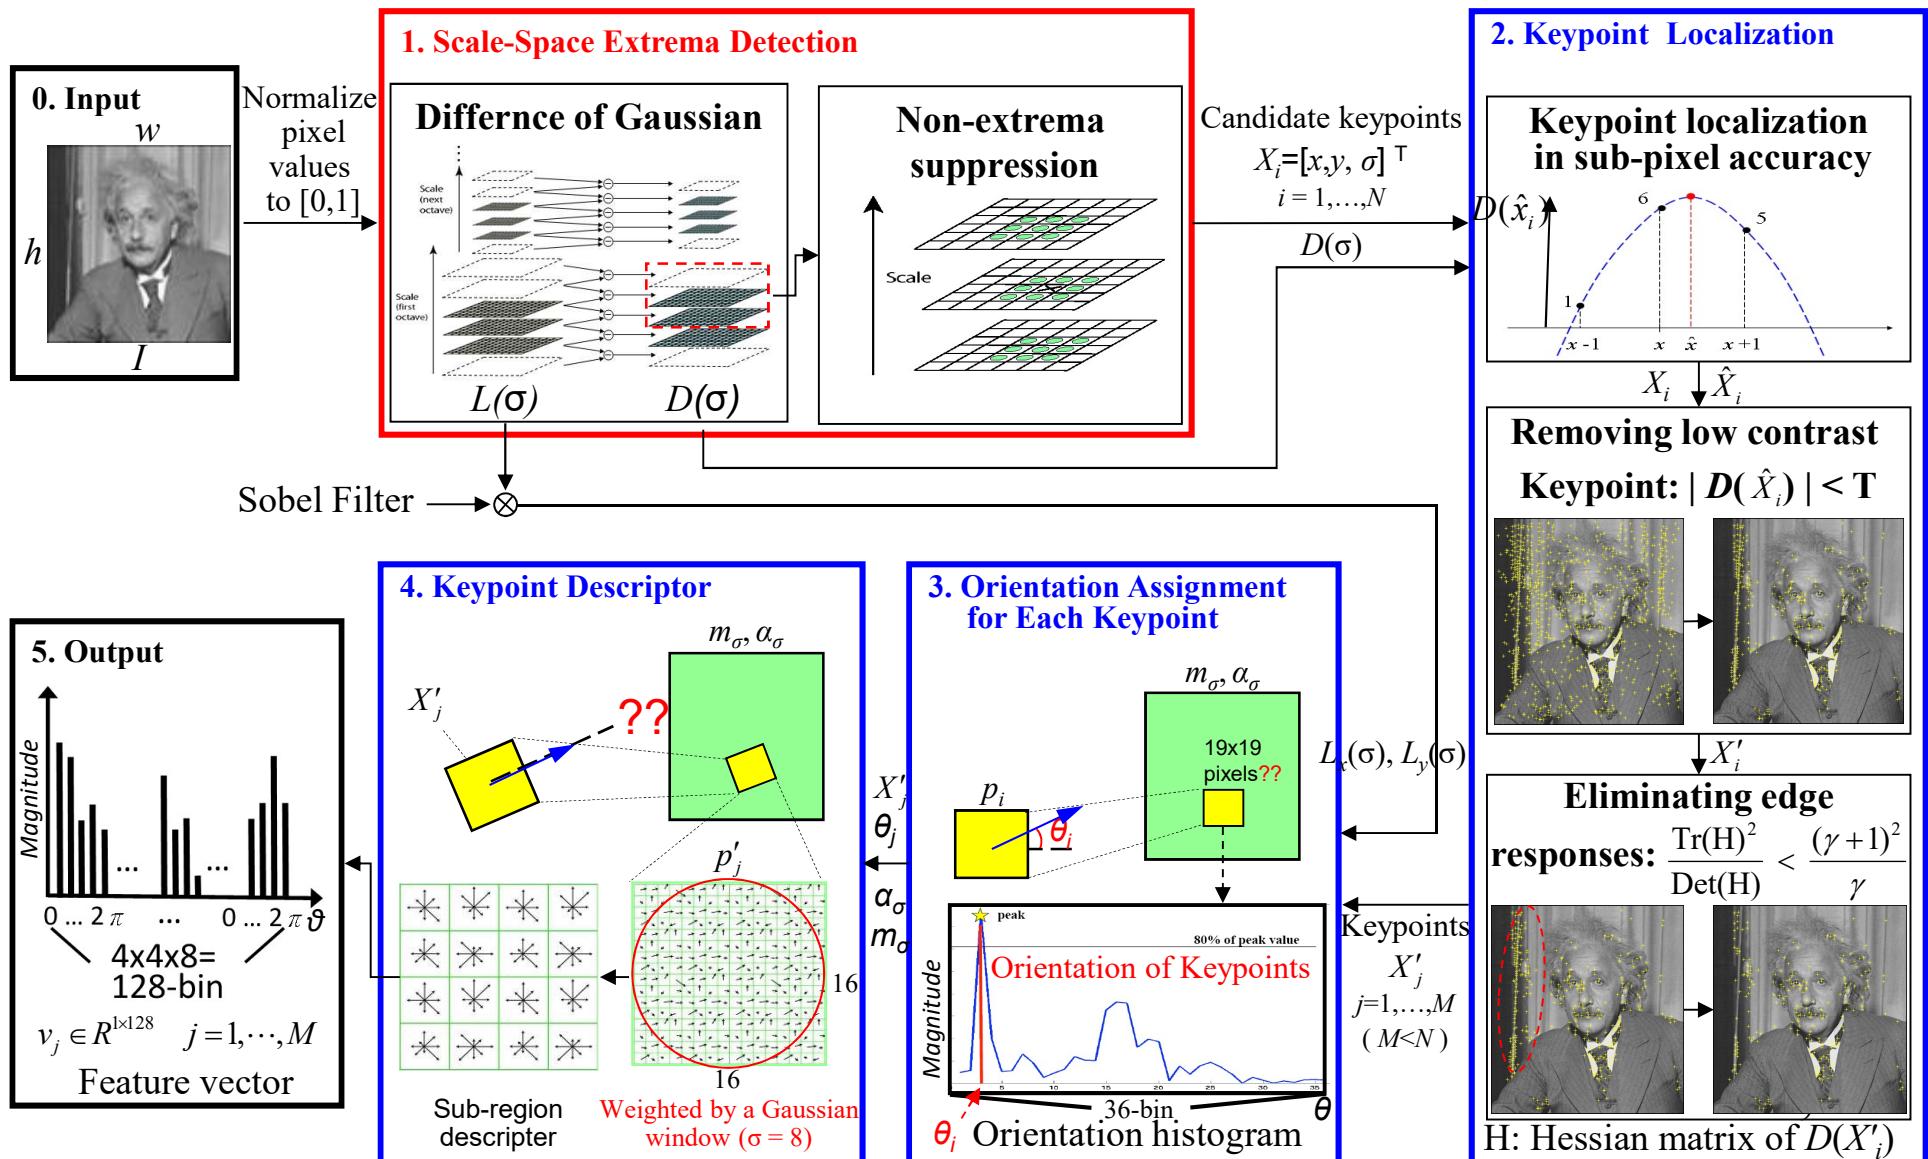
$$L(\sigma) = G(\sigma) * I$$

$$D(\sigma) = (G(k^{(n-1)}\sigma) - G(k^n\sigma)) * I$$

(Scale space)
(DoG images)

α_σ (Gradient orientation of $L(\sigma)$)
 m_σ (Gradient magnitude of $L(\sigma)$)
 X'_i (Accurate keypoint localization)
 θ_j (Orientation of i -th keypoint)

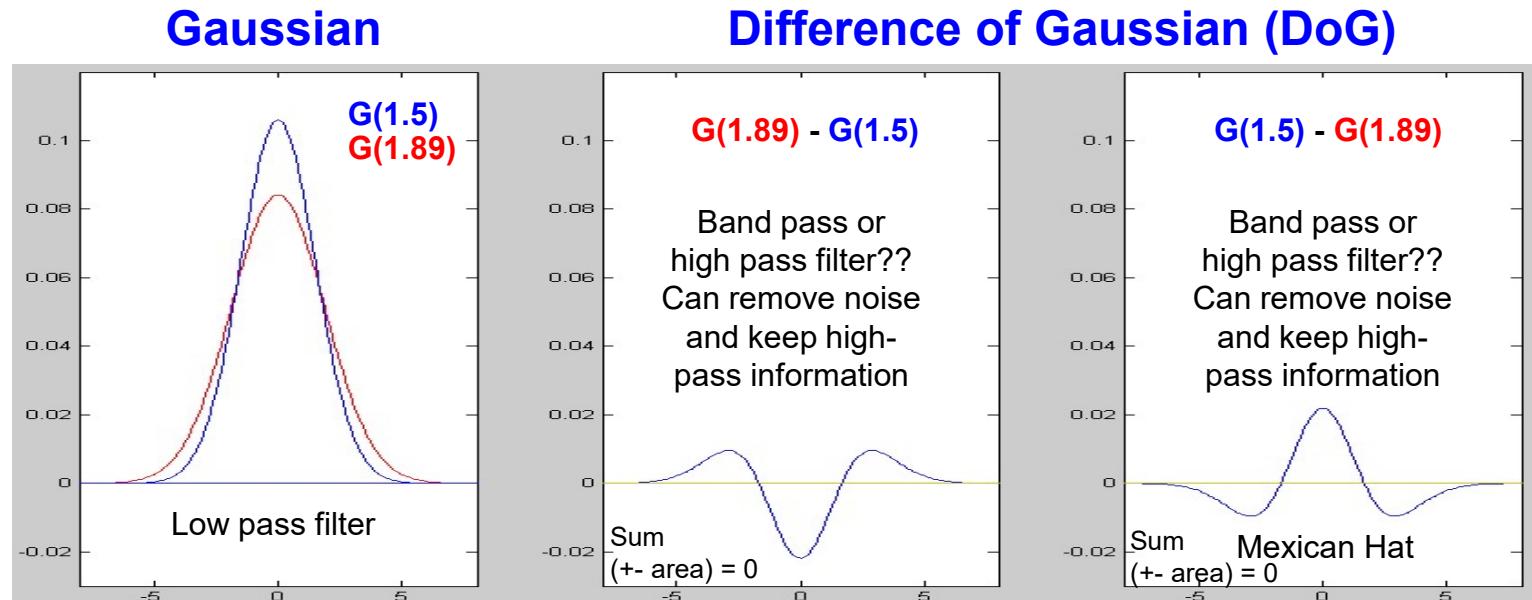
System Flowchart



3. Detection of Scale-Space Extrema

◆ Difference of Gaussian (DoG)

- For **scale invariance**, search for stable features across all possible scales using a continuous function of scale, such as Gaussian.
- SIFT uses **DoG filter** for scale space because it is efficient and as stable as scale-normalized Laplacian of Gaussian (LOG).



低通與高通濾波器

- 高頻(high-frequency components)：短距離內灰階值變化值大，intensity gradient 大。(例：texture、邊緣或雜訊)。
- 低頻(low-frequency components)：短距離內灰階值變化值小，intensity gradient 小。(例：textureless)
- 高通濾波器 (high-pass filter)：保持高頻部分，減少低頻部分。(係數總和為0)(例如 Sobel 或 Canny)
- 低通濾波器 (low-pass filter)：保持低頻部分，減少高頻部分。Smooth (係數總和為1)(例如 Gaussian Filter)

Example

$$\begin{array}{ll} \text{high-pass filter} & \text{low-pass filter} \\ \begin{bmatrix} 1 & -2 & 1 \\ -2 & 4 & -2 \\ 1 & -2 & 1 \end{bmatrix} & \frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix} \end{array}$$

Q&A. DoG vs. LoG (1/2)

- Difference of Gaussian operation provides an **efficient approximation** to the Laplacian.
- In reality, Laplacian has better performance than DoG but slow.

Q&A. DoG vs. LoG (2/2)

✓ LoG (Laplacian of Gaussian)

- Laplacian operator: $\nabla^2 = \frac{\partial^2}{\partial x^2} + \frac{\partial^2}{\partial y^2}$
- 2-D Gaussian function: $G(x, y) = e^{-\frac{x^2+y^2}{2\sigma^2}}$

$$g(x, y) = [\nabla^2 G(x, y)] \star f(x, y)$$

$$g(x, y) = \nabla^2 [G(x, y) \star f(x, y)]$$

finding the **zero crossings** of $g(x, y)$ to determine the locations of edges in $f(x, y)$.

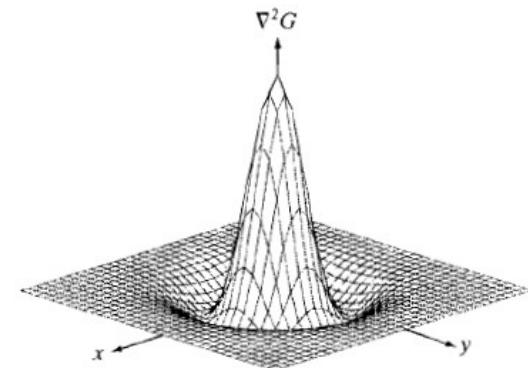
$$\nabla^2 G(x, y) = \frac{\partial^2 G(x, y)}{\partial x^2} + \frac{\partial^2 G(x, y)}{\partial y^2}$$

$$= \frac{\partial}{\partial x} \left[\frac{-x}{\sigma^2} e^{-\frac{x^2+y^2}{2\sigma^2}} \right] + \frac{\partial}{\partial y} \left[\frac{-y}{\sigma^2} e^{-\frac{x^2+y^2}{2\sigma^2}} \right]$$

$$= \left[\frac{x^2}{\sigma^4} - \frac{1}{\sigma^2} \right] e^{-\frac{x^2+y^2}{2\sigma^2}} + \left[\frac{y^2}{\sigma^4} - \frac{1}{\sigma^2} \right] e^{-\frac{x^2+y^2}{2\sigma^2}}$$

$$= \left[\frac{x^2 + y^2 - 2\sigma^2}{\sigma^4} \right] e^{-\frac{x^2+y^2}{2\sigma^2}}$$

$g(x, y)$: LoG image.
 $f(x, y)$: Input image.
 \star : Convolution.



LoG的三維模型

	-2	-1	0	1	2
-2	0	0	-1	0	0
-1	0	-1	-2	-1	0
0	-1	-2	16	-2	-1
1	0	-1	-2	-1	0
2	0	0	-1	0	0

|o=2

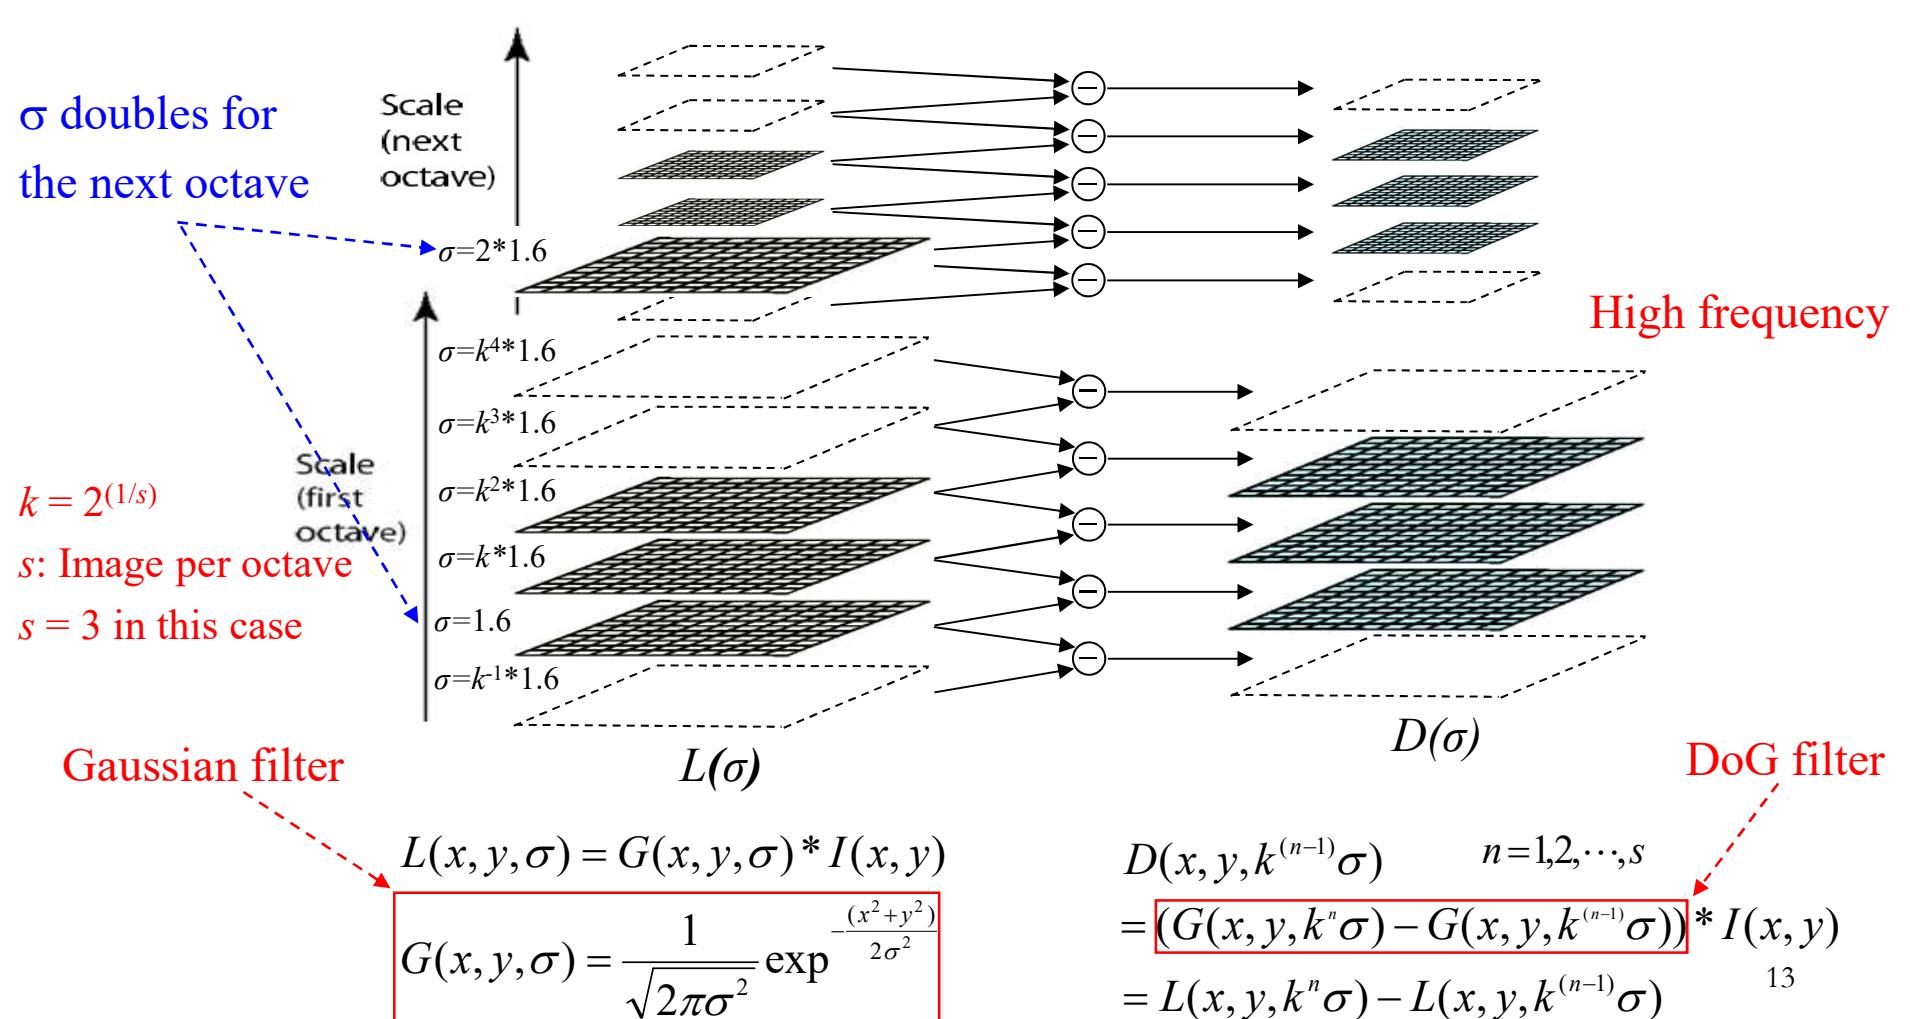
近似的5*5 mask

s is incorrect!!

3.1 Difference of Gaussian (DoG) filter (1/2)

Dividing into octave is for efficiency only

: Until image size < 32x32



3.1 Difference of Gaussian (DoG) filter (1/2)

$$L(x, y, \sigma) = G(x, y, \sigma) * I(x, y)$$

$$D(x, y, k^{(n-1)}\sigma) = (G(x, y, k^n\sigma) - G(x, y, k^{(n-1)}\sigma)) * I(x, y)$$

$n=1, 2, \dots, s+2$

$$= L(x, y, k^n\sigma) - L(x, y, k^{(n-1)}\sigma)$$

$$G(x, y, \sigma) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp^{-\frac{(x^2+y^2)}{2\sigma^2}}$$

$G(x, y, \sigma)$: Gaussian filter.
 $I(x, y)$: Input Image.

How many images in each octave /layer?

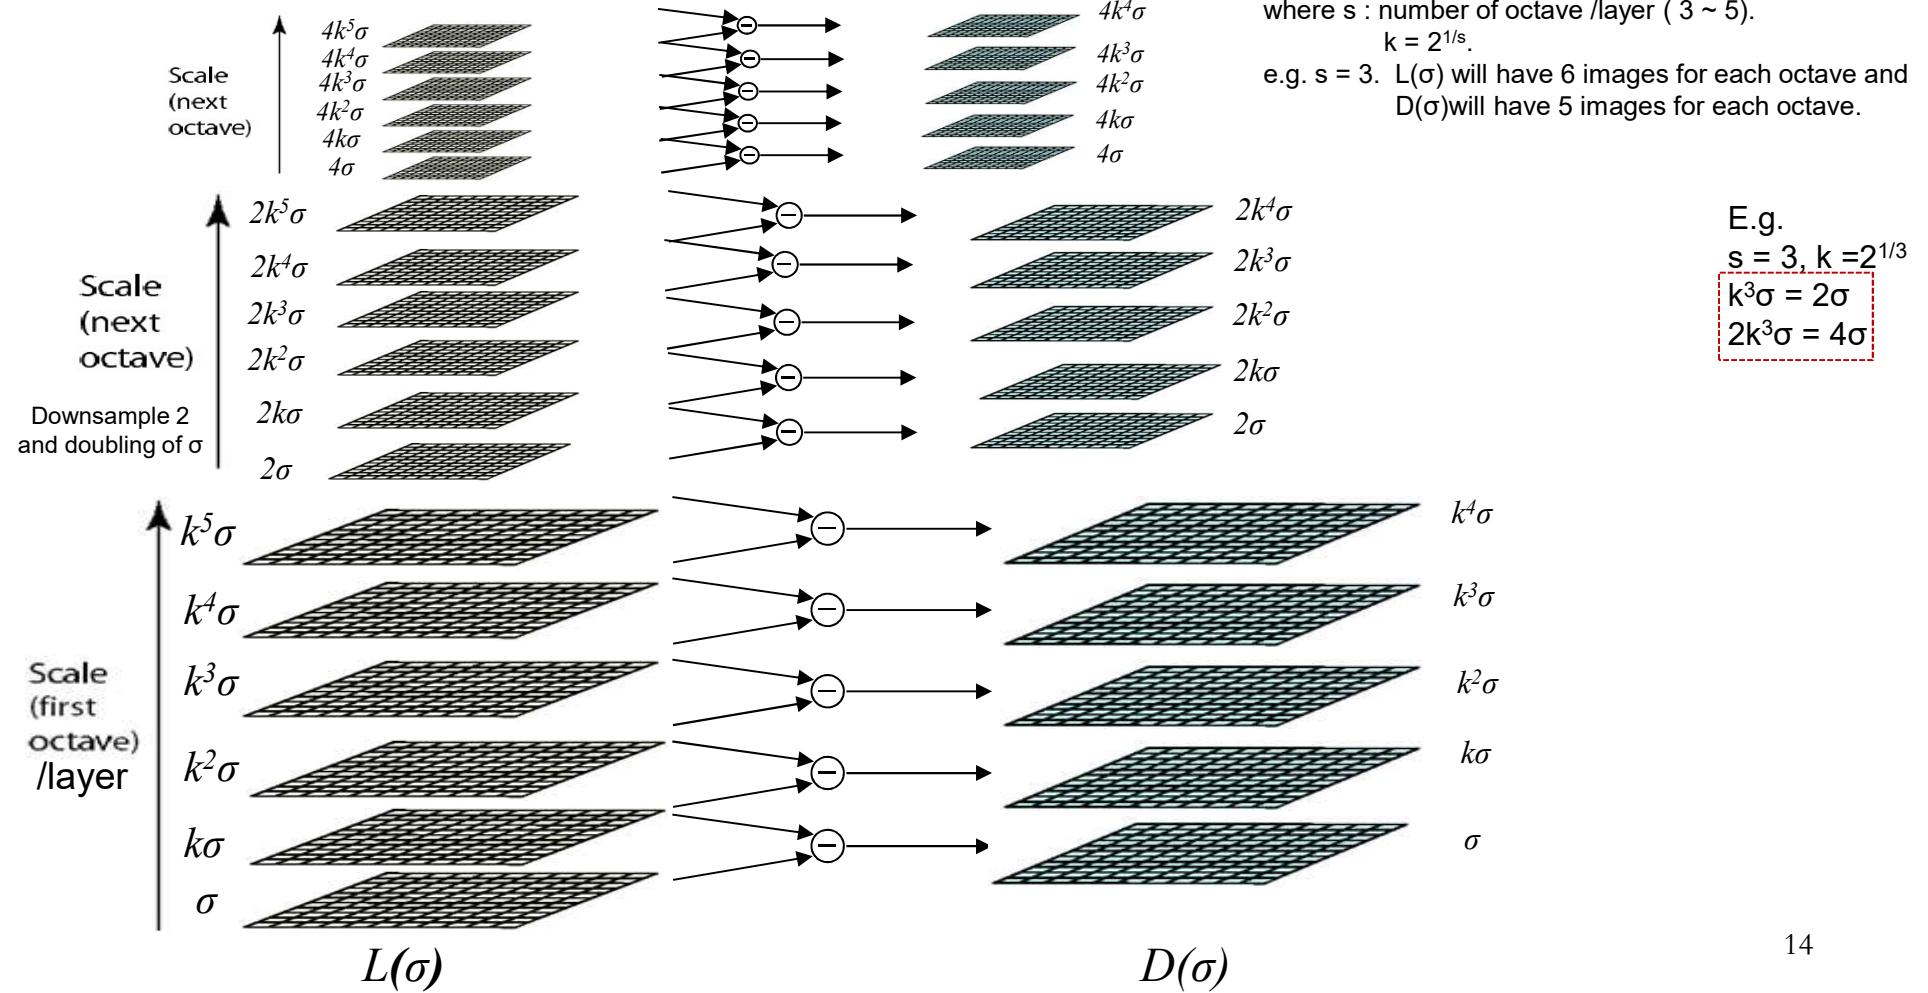
$L(\sigma)$ will have $s+3$ images for each octave.

$D(\sigma)$ will have $s+2$ images for each octave.

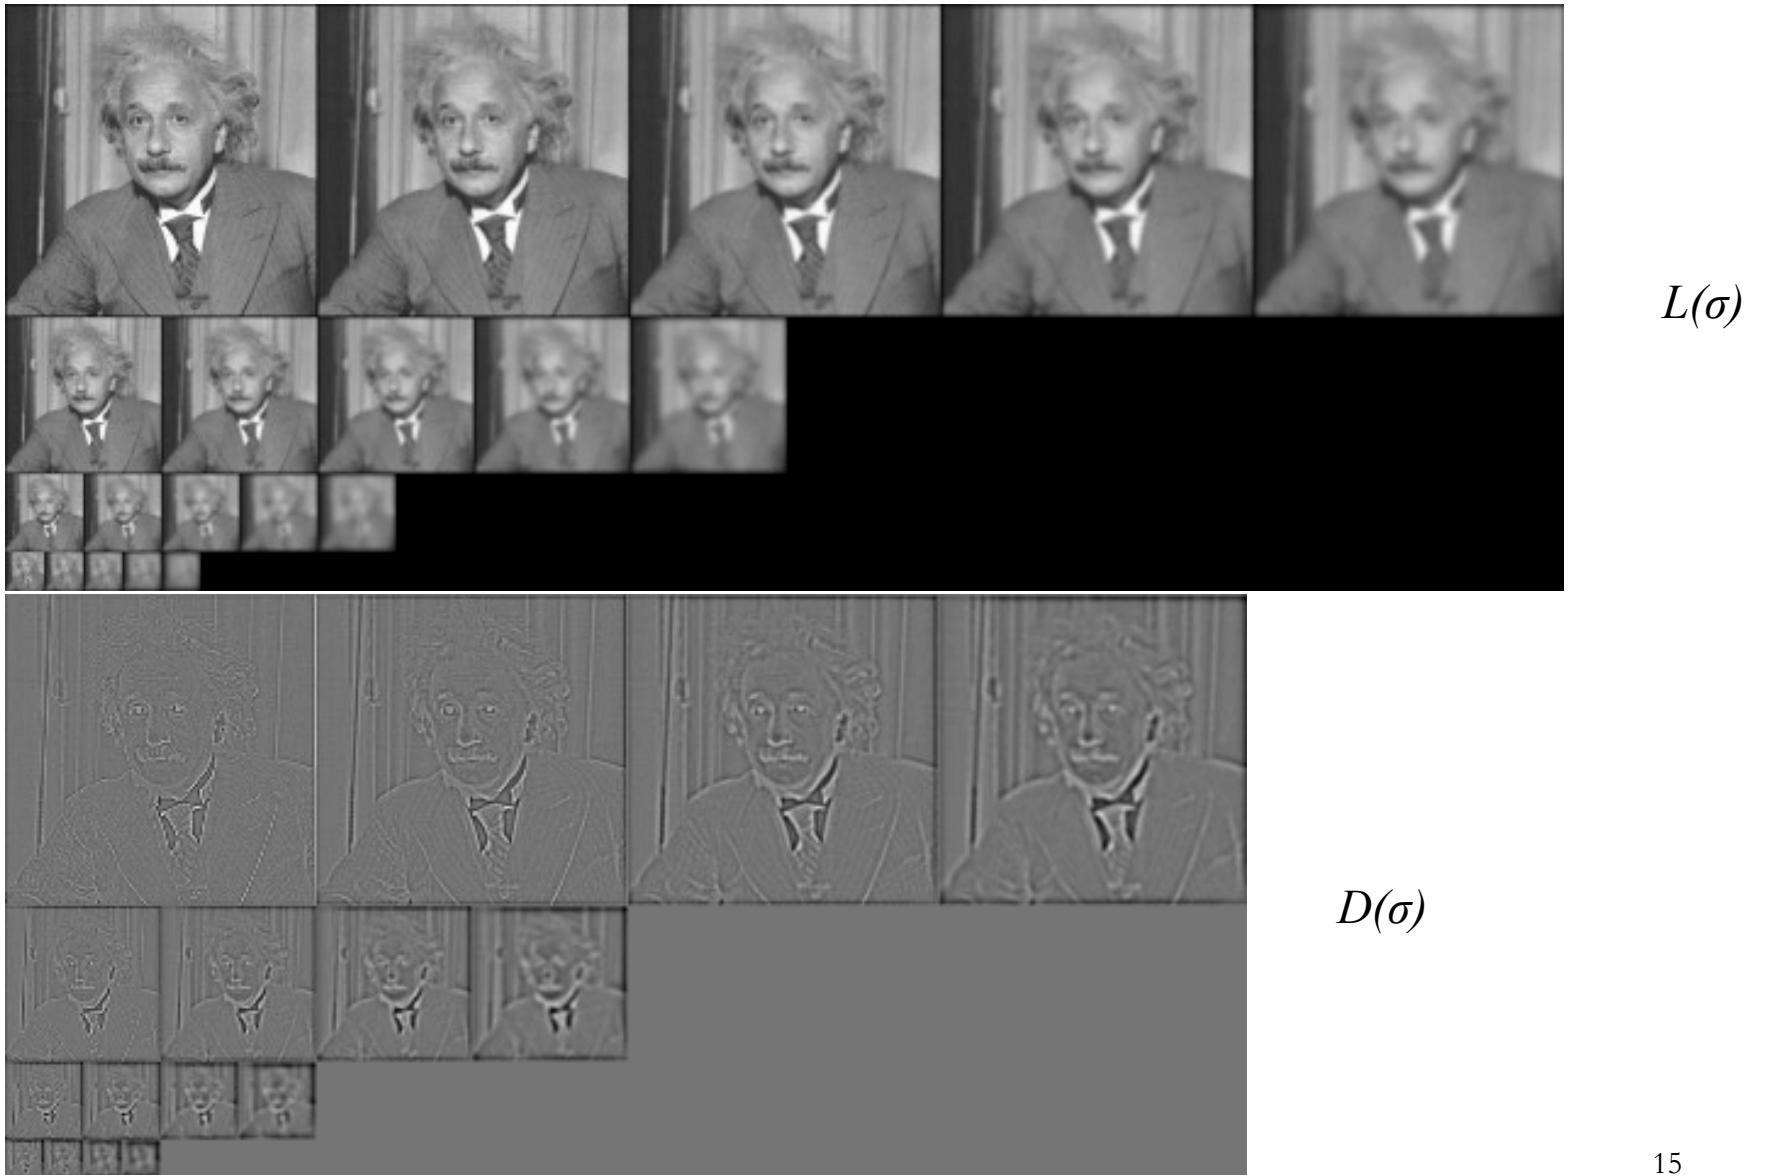
where s : number of octave /layer ($3 \sim 5$).

$$k = 2^{1/s}$$

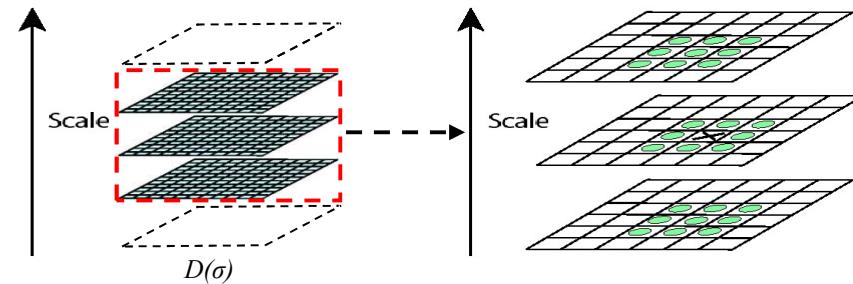
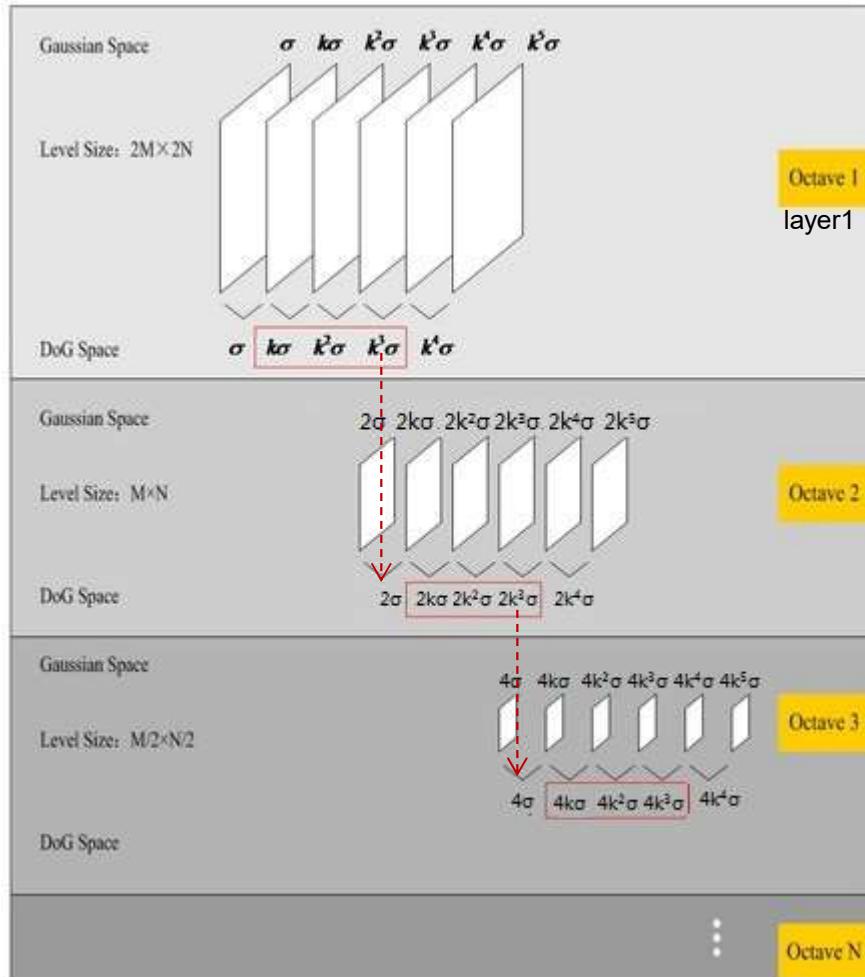
e.g. $s = 3$. $L(\sigma)$ will have 6 images for each octave and $D(\sigma)$ will have 5 images for each octave.



3.1 Difference of Gaussian (DoG) filter (2/2)



3.2 Non-maxima/minima suppression

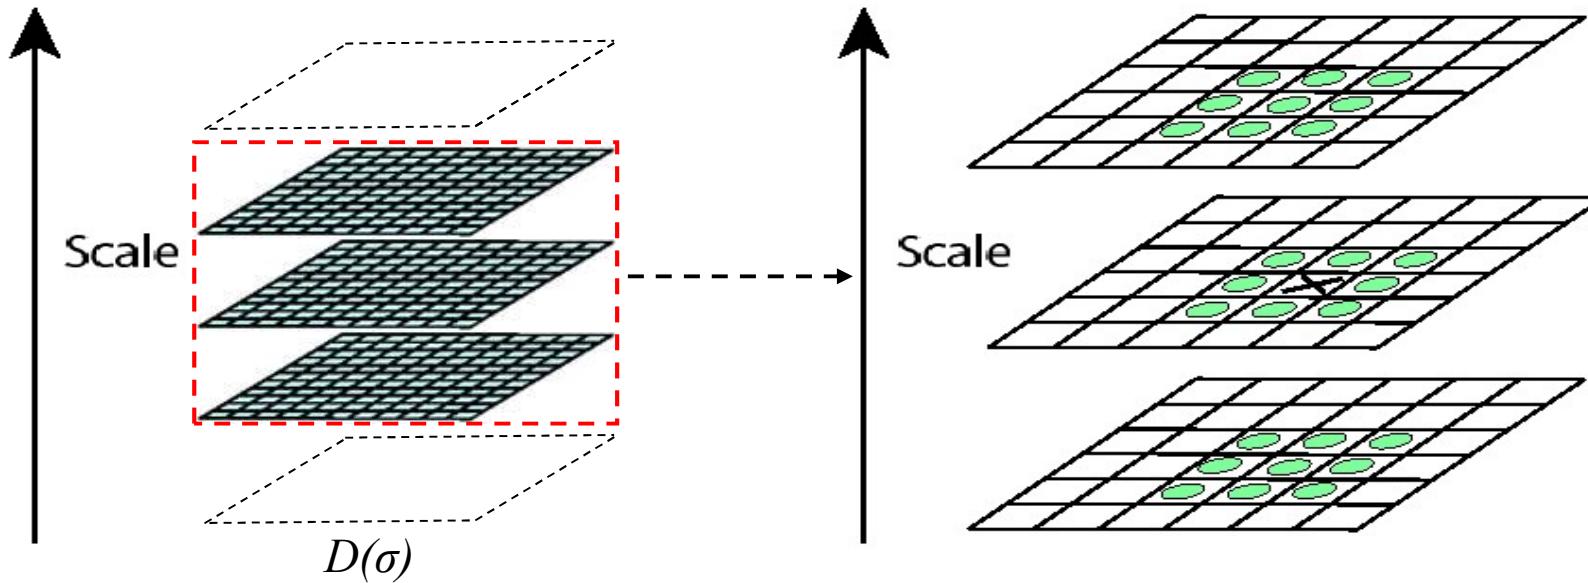


- $X_i = (x, y, \sigma)^T$ is selected as a candidate keypoint if it is **larger** or **smaller** than all 26 neighbors.
- Local maximum/minimum: After local extrema detection, we find extrema point as: $X_i = (x, y, \sigma)^T$

E.g.

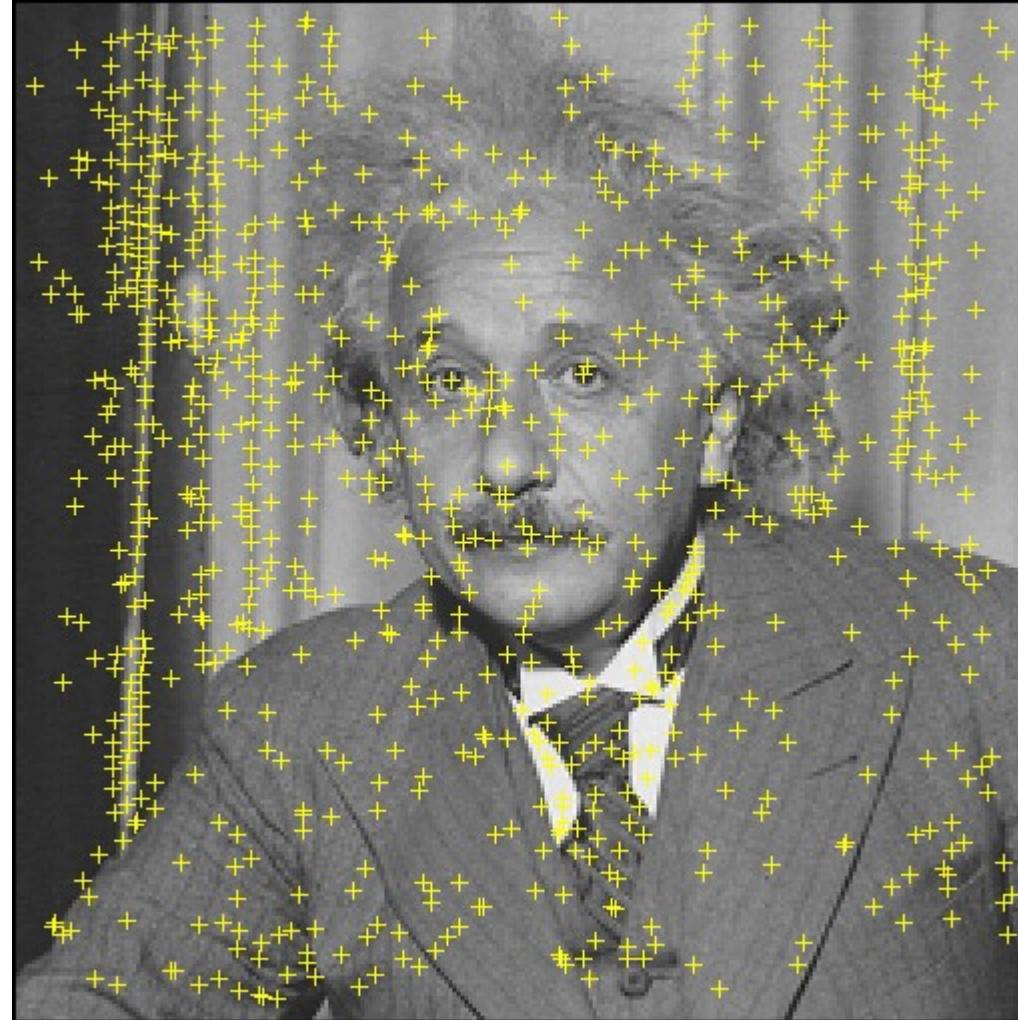
$$\begin{aligned}s &= 3, k = 2^{1/3} \\ k^3\sigma &= 2\sigma \\ 2k^3\sigma &= 4\sigma\end{aligned}$$

3.2 Non-maxima/minima suppression (1/2) ??



- x is selected as a candidate **keypoint** if it is **larger** or **smaller** than all 26 neighbors.
- Local maximum/minimum
- Q? So does this candidate keypoint exist any 8-connected component keypoints? Probably Not

3.2 Non-maxima/minima suppression (2/2)



Extrema in D (Candidate keypoints)

Content

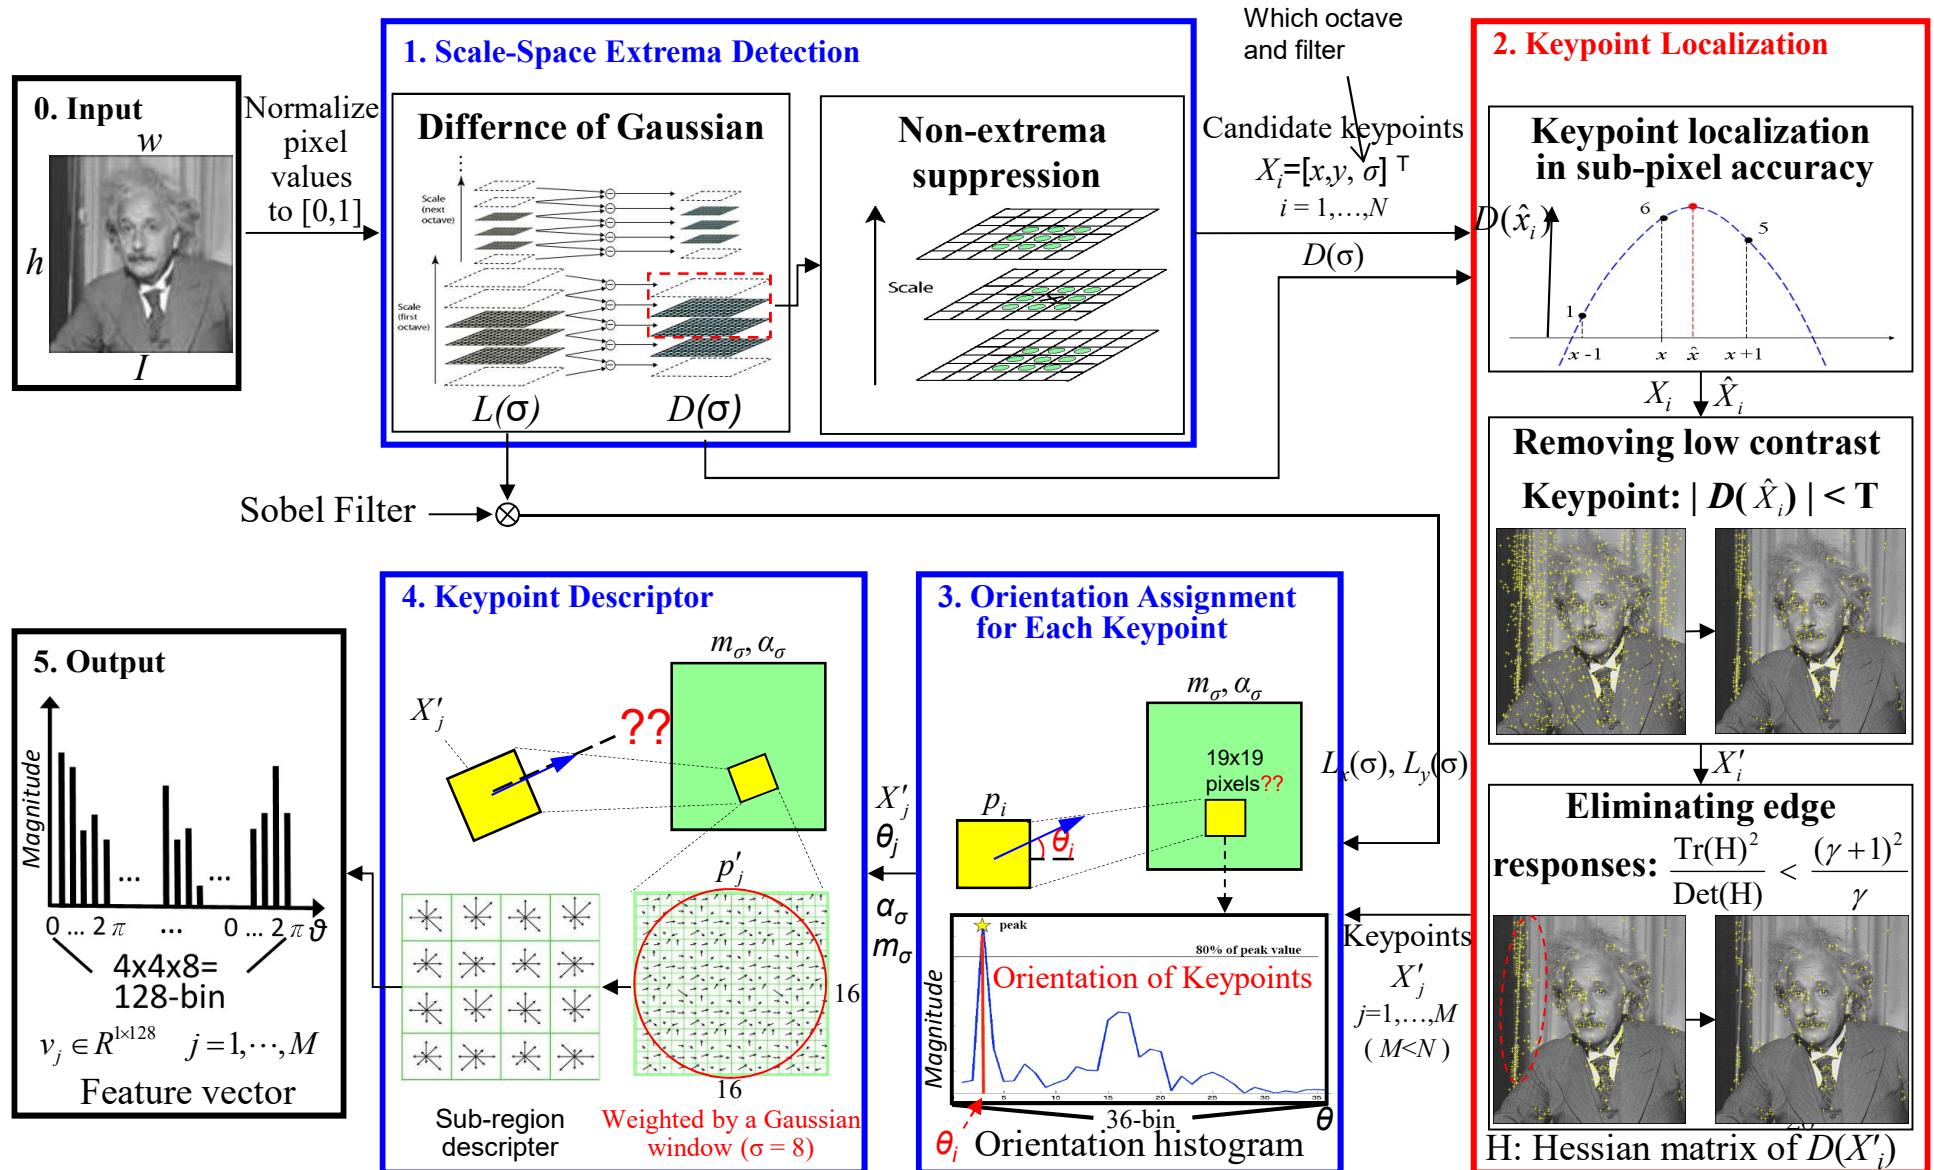
1. Motivation
2. System Flowchart
3. Detection of Scale-Space Extrema
 - 3.1 Difference of Gaussian (DoG) filter
 - 3.2 Non-maxima/minima suppression
4. Keypoint Localization
 - 4.1 Keypoint localization in sub-pixel accuracy
 - 4.2 Remove low contrast keypoints
 - 4.3 Eliminating edge responses
5. Orientation Assignment
6. Keypoint Descriptor
7. Parameter Selection

$$L(\sigma) = G(\sigma) * I$$

$$D(\sigma) = (G(k^{(n-1)}\sigma) - G(k^n\sigma)) * I$$

(Scale space)	α_σ	(Gradient orientation of $L(\sigma)$)
(DoG images)	m_σ	(Gradient magnitude of $L(\sigma)$)
	X'_i	(Accurate keypoint localization)
	θ_j	(Orientation of i -th keypoint)

System Flowchart



4. Accurate Keypoint Localization 1/7

Once a keypoint candidate has been found by comparing a pixel to its neighbors, the next step is to perform a detailed fit to the nearby data for location, scale, and ratio of principal curvatures. This information allows points to be rejected that have low contrast (and are therefore sensitive to noise) or are poorly localized along an edge.

The initial implementation of this approach (Lowe, 1999) simply located keypoints at the location and scale of the central sample point. However, recently Brown has developed a method (Brown and Lowe, 2002) for fitting a 3D quadratic function to the local sample points to determine the interpolated location of the maximum, and his experiments showed that this provides a substantial improvement to matching and stability. His approach uses the Taylor expansion (up to the quadratic terms) of the scale-space function, $D(x, y, \sigma)$, shifted so that the origin is at the sample point:

$$D(\mathbf{x}) = D + \frac{\partial D^T}{\partial \mathbf{x}} \mathbf{x} + \frac{1}{2} \mathbf{x}^T \frac{\partial^2 D}{\partial \mathbf{x}^2} \mathbf{x} \quad (2)$$

where D and its derivatives are evaluated at the sample point and $\mathbf{x} = (x, y, \sigma)^T$ is the offset from this point.

The location of the extremum, $\hat{\mathbf{x}}$, is determined by taking the derivative of this function with respect to \mathbf{x} and setting it to zero, giving

$$\hat{\mathbf{x}} = -\frac{\partial^2 D^{-1}}{\partial \mathbf{x}^2} \frac{\partial D}{\partial \mathbf{x}}. \quad (3)$$

As suggested by Brown, the Hessian and derivative of D are approximated by using differences of neighboring sample points. The resulting 3×3 linear system can be solved with minimal cost. If the offset $\hat{\mathbf{x}}$ is larger than 0.5 in any dimension, then it means that the extremum lies closer to a different sample point. In this case, the sample point is changed and the interpolation performed instead about that point. The final offset $\hat{\mathbf{x}}$ is added to the location of its sample point to get the interpolated estimate for the location of the extremum.

The function value at the extremum, $D(\hat{\mathbf{x}})$, is useful for rejecting unstable extrema with low contrast. This can be obtained by substituting Eqs (3) into (2), giving

$$D(\hat{\mathbf{x}}) = D + \frac{1}{2} \frac{\partial D^T}{\partial \mathbf{x}} \hat{\mathbf{x}}.$$

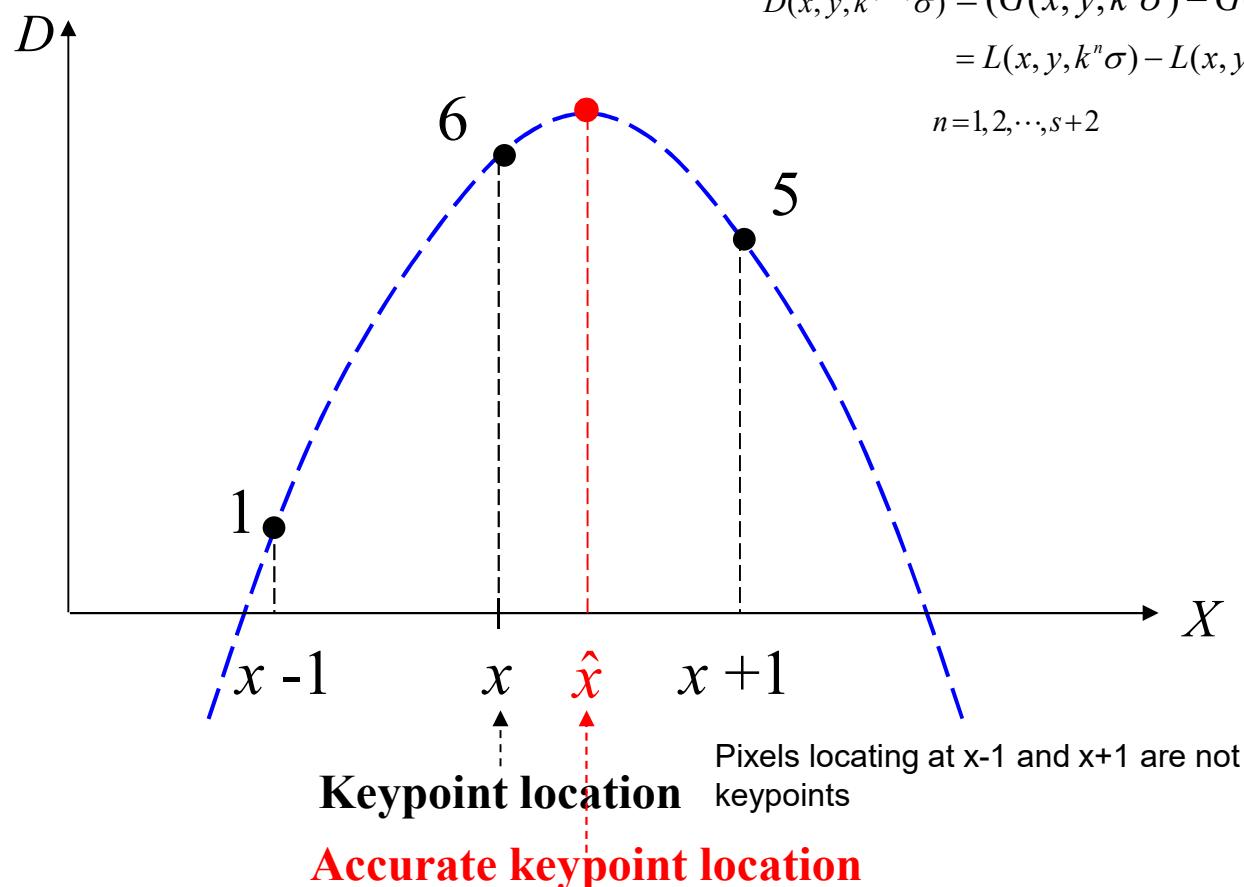
For the experiments in this paper, all extrema with a value of $D(\hat{\mathbf{x}})$ less than 0.03 were discarded (as before, we assume image pixel values in the range [0, 1]).

4.1 Accurate keypoint localization (2/7)

- Perform a detailed fit to the nearby data around the keypoint for accurate location by **Taylor expansion** (up to the quadratic terms).

$$\begin{aligned}L(x, y, \sigma) &= G(x, y, \sigma) * I(x, y) \\D(x, y, k^{(n-1)}\sigma) &= (G(x, y, k^n\sigma) - G(x, y, k^{(n-1)}\sigma)) * I(x, y) \\&= L(x, y, k^n\sigma) - L(x, y, k^{(n-1)}\sigma)\end{aligned}$$

$n=1, 2, \dots, s+2$

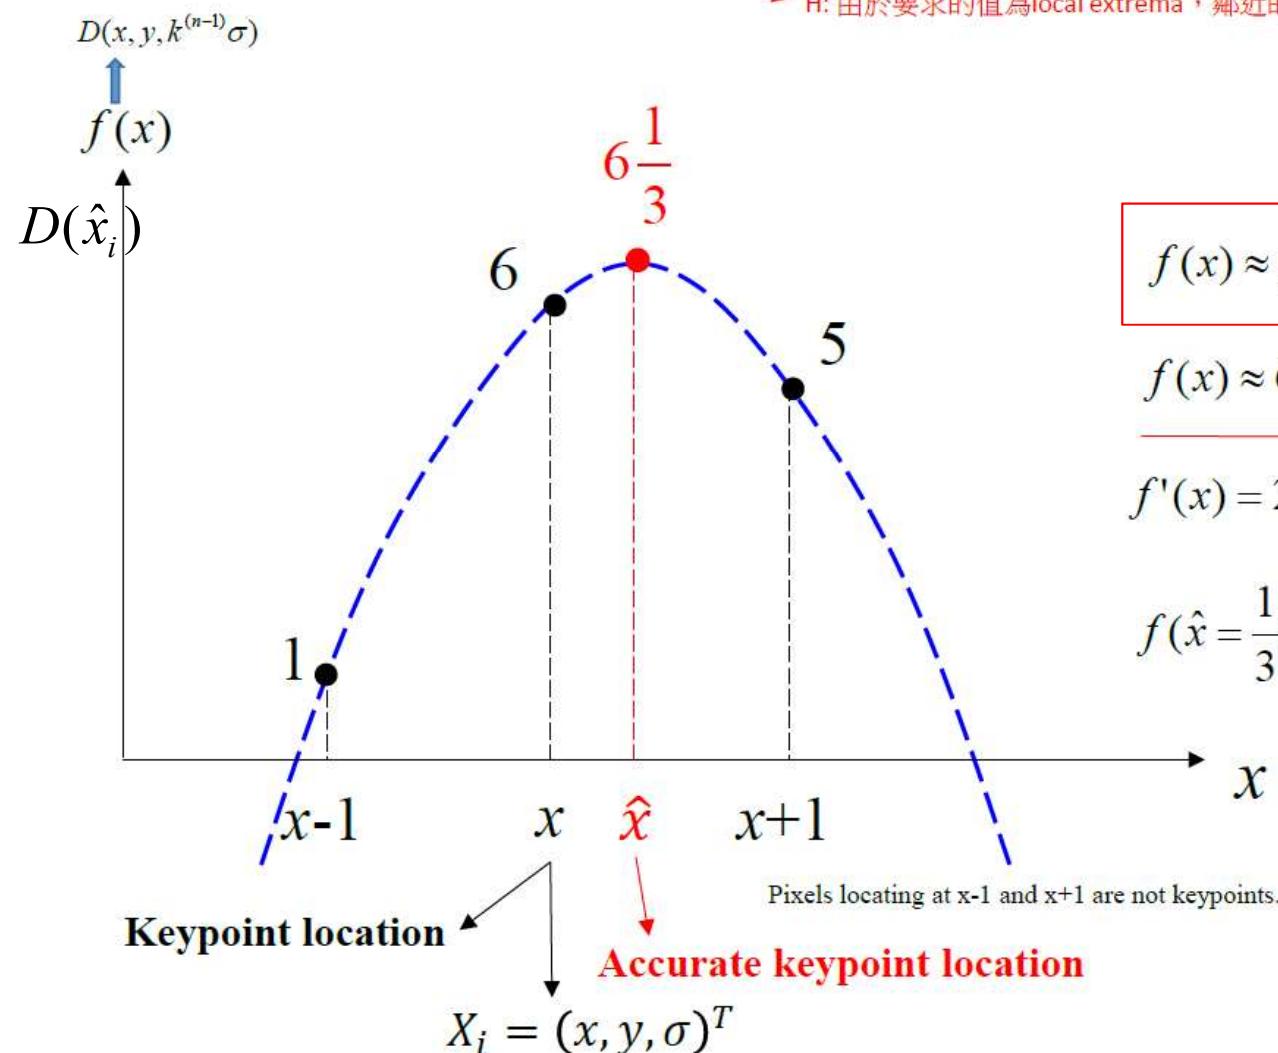


Nonlinear Interpolation!!

4.1 Accurate keypoint localization (3/7)

2.1) Accurate Feature Localization by Taylor expansion J: Ideal solution but slow

Perform a detailed fit to the nearby data around the keypoint for accurate location by Taylor expansion (up to the quadratic terms). Why quadratic, not linear interpolation??



H: 由於要求的值為local extrema，鄰近的點分布會是曲線，因此要用二階來求方程式。

$$f(x) \approx f(0) + f'(0)x + \frac{f''(0)}{2}x^2$$

$$f(x) \approx 6 + 2x + \frac{-6}{2}x^2 = 6 + 2x - 3x^2$$

$$f'(x) = 2 - 6x = 0, \text{ so } x = \frac{1}{3} = \hat{x}$$

(Next page)

$$f(\hat{x} = \frac{1}{3}) = 6 + 2 \cdot \frac{1}{3} - 3 \cdot \left(\frac{1}{3}\right)^2 = 6 \frac{1}{3}$$

4.1 Accurate keypoint localization (4/7)

求 $f(x)$

$$f(x) \approx f(0) + f'(0)x + \frac{f''(0)}{2} x^2$$

3 unknown parameters?
Need 3 1D points => 3 equations

$$\text{if } x = 0, \quad f(x) = f(0) + f'(0) * 0 + \frac{f''(0)}{2} * 0^2 = 6$$

$$if \ x = -1, \ f(x) = f(0) + f'(0) * (-1) + \frac{f''(0)}{2} * (-1)^2$$

$$= 6 - f'(0) + \frac{f''(0)}{2} = \boxed{1}$$

$$f'(0) - \frac{f''(0)}{2} = 5 \quad \dots \dots \dots (1)$$

$$if \ x = 1, \quad f(x) = f(0) + f'(0) * 1 + \frac{f''(0)}{2} * 1^2$$

$$= 6 + f'(0) + \frac{f''(0)}{2} = 5$$

$$f'(0) + \frac{f''(0)}{2} = -1 \quad \dots \dots \dots (2)$$

$$(1)+(2) \quad 2f'(0) = 4 \quad f'(0) = 2$$

$$\text{代回(2)} \quad \frac{f''(0)}{2} = -3 \quad f''(0) = -6$$

$$f(x) = 6 + 2x - 3x^2$$

4.1 Accurate keypoint localization (5/7)

- Taylor series of several variables

$$T(x_1, \dots, x_d) = \sum_{n_1=0}^{\infty} \dots \sum_{n_d=0}^{\infty} \frac{\partial^{n_1}}{\partial x_1^{n_1}} \dots \frac{\partial^{n_d}}{\partial x_d^{n_d}} \frac{f(a_1, \dots, a_d)}{n_1! \dots n_d!} (x_1 - a_1)^{n_1} \dots (x_d - a_d)^{n_d}$$

- Taylor expansion in matrix form, \mathbf{x} is a $n \times 1$ vector, f maps \mathbf{x} to a scalar.

$$f(\mathbf{x}) = f + \frac{\partial f}{\partial \mathbf{x}}^T \mathbf{x} + \frac{1}{2} \mathbf{x}^T \frac{\partial^2 f}{\partial \mathbf{x}^2} \mathbf{x}$$

gradient
$$\begin{pmatrix} \frac{\partial f}{\partial x_1} \\ \frac{\partial f}{\partial x_2} \\ \vdots \\ \frac{\partial f}{\partial x_n} \end{pmatrix}$$

↓ ↓

Hessian
$$\begin{pmatrix} \frac{\partial^2 f}{\partial x_1^2} & \frac{\partial^2 f}{\partial x_1 \partial x_2} & \dots & \frac{\partial^2 f}{\partial x_1 \partial x_n} \\ \frac{\partial^2 f}{\partial x_2 \partial x_1} & \frac{\partial^2 f}{\partial x_2^2} & \dots & \frac{\partial^2 f}{\partial x_2 \partial x_n} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial^2 f}{\partial x_n \partial x_1} & \frac{\partial^2 f}{\partial x_n \partial x_2} & \dots & \frac{\partial^2 f}{\partial x_n^2} \end{pmatrix}$$

4.1 Accurate keypoint localization (6/7) ??

DoG

- Fit $D(x, y, \sigma)$ quadratic function for sub-pixel extrema.

$$D(\mathbf{x}) \approx D + \frac{\partial D^T}{\partial \mathbf{x}} \mathbf{x} + \frac{1}{2} \mathbf{x}^T \frac{\partial^2 D}{\partial \mathbf{x}^2} \mathbf{x}$$

- \mathbf{x} is a 3×1 vector $[x \ y \ \sigma]^T$

?? $\frac{\partial D(\mathbf{x})}{\partial \mathbf{x}} = \frac{\partial D^T}{\partial \mathbf{x}} + \frac{1}{2} \left(\frac{\partial^2 D}{\partial \mathbf{x}^2} + \frac{\partial^2 D^T}{\partial \mathbf{x}^2} \right) \mathbf{x} = \frac{\partial D^T}{\partial \mathbf{x}} + \frac{\partial^2 D}{\partial \mathbf{x}^2} \mathbf{x} = 0$

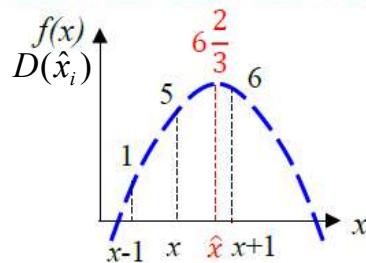
Sub-pixel
extrema
position

$$\hat{\mathbf{x}} = -\frac{\partial^2 D}{\partial \mathbf{x}^2}^{-1} \frac{\partial D}{\partial \mathbf{x}}$$

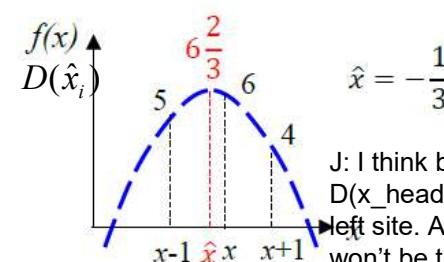
J: Similar to optical flow

- If the offset \hat{x} is larger than 0.5 in any dimension, then it means that the extrema lies closer to a different sample point. In this case, **the sample point is changed and the interpolation performed instead about that point**. The final offset \hat{x} is added to the location of its sample point to get the interpolated estimate for the location of the extrema.

Example



$\hat{x} = \frac{2}{3} > 0.5$
Sample point change to $x+1$



J: I think both x_{head} value and $D(x_{head})$ at right side are not as left site. And the quadratic curve won't be the same either

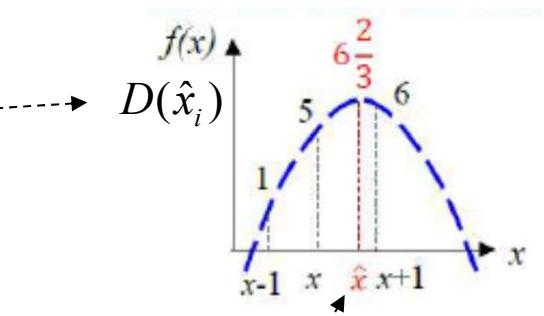
4.2 Remove low contrast keypoints (7/7)

2.2) Removing Low Contrast Feature

- Throw out low contrast features which is $|D(\hat{x}_i)| < 0.03$

$$D(\hat{x}_i) = D + \frac{1}{2} \frac{\partial D^T}{\partial x_i} \hat{x}_i$$

Original value $\frac{1}{2} \frac{\partial D^T}{\partial x_i}$ additional value



derive: $D(\hat{\mathbf{x}}) \approx D + \frac{\partial D^T}{\partial \mathbf{x}} \hat{\mathbf{x}} + \frac{1}{2} \hat{\mathbf{x}}^T \frac{\partial^2 D}{\partial \mathbf{x}^2} \hat{\mathbf{x}}$

$$= D + \frac{\partial D^T}{\partial \mathbf{x}} \hat{\mathbf{x}} + \frac{1}{2} \left(-\frac{\partial^2 D^{-1}}{\partial \mathbf{x}^2} \frac{\partial D}{\partial \mathbf{x}} \right)^T \frac{\partial^2 D}{\partial \mathbf{x}^2} \left(-\frac{\partial^2 D^{-1}}{\partial \mathbf{x}^2} \frac{\partial D}{\partial \mathbf{x}} \right)$$

$$= D + \frac{\partial D^T}{\partial \mathbf{x}} \hat{\mathbf{x}} + \frac{1}{2} \frac{\partial D^T}{\partial \mathbf{x}} \cancel{\frac{\partial^2 D^{-1}}{\partial \mathbf{x}^2}} \cancel{\frac{\partial^2 D}{\partial \mathbf{x}^2}} \frac{\partial^2 D^{-1}}{\partial \mathbf{x}^2} \frac{\partial D}{\partial \mathbf{x}}$$

$$= D + \frac{\partial D^T}{\partial \mathbf{x}} \hat{\mathbf{x}} + \frac{1}{2} \frac{\partial D^T}{\partial \mathbf{x}} \underline{\frac{\partial^2 D^{-1}}{\partial \mathbf{x}^2}} \frac{\partial D}{\partial \mathbf{x}}$$

$$= D + \frac{\partial D^T}{\partial \mathbf{x}} \hat{\mathbf{x}} + \frac{1}{2} \frac{\partial D^T}{\partial \mathbf{x}} (-\hat{\mathbf{x}})$$

$$= D + \frac{1}{2} \frac{\partial D^T}{\partial \mathbf{x}} \hat{\mathbf{x}}$$

$$\hat{\mathbf{x}} = -\frac{\partial^2 D^{-1}}{\partial \mathbf{x}^2} \frac{\partial D}{\partial \mathbf{x}}$$

J: If Alfa and Beta are eigenvalues, then probably the equations are trying to solve both eigenvalues ??

4.3 Eliminating edge responses (1/2)

2.3) Eliminating Edge Responses

- Compute 2×2 Hessian matrix at the location and scale of the feature:

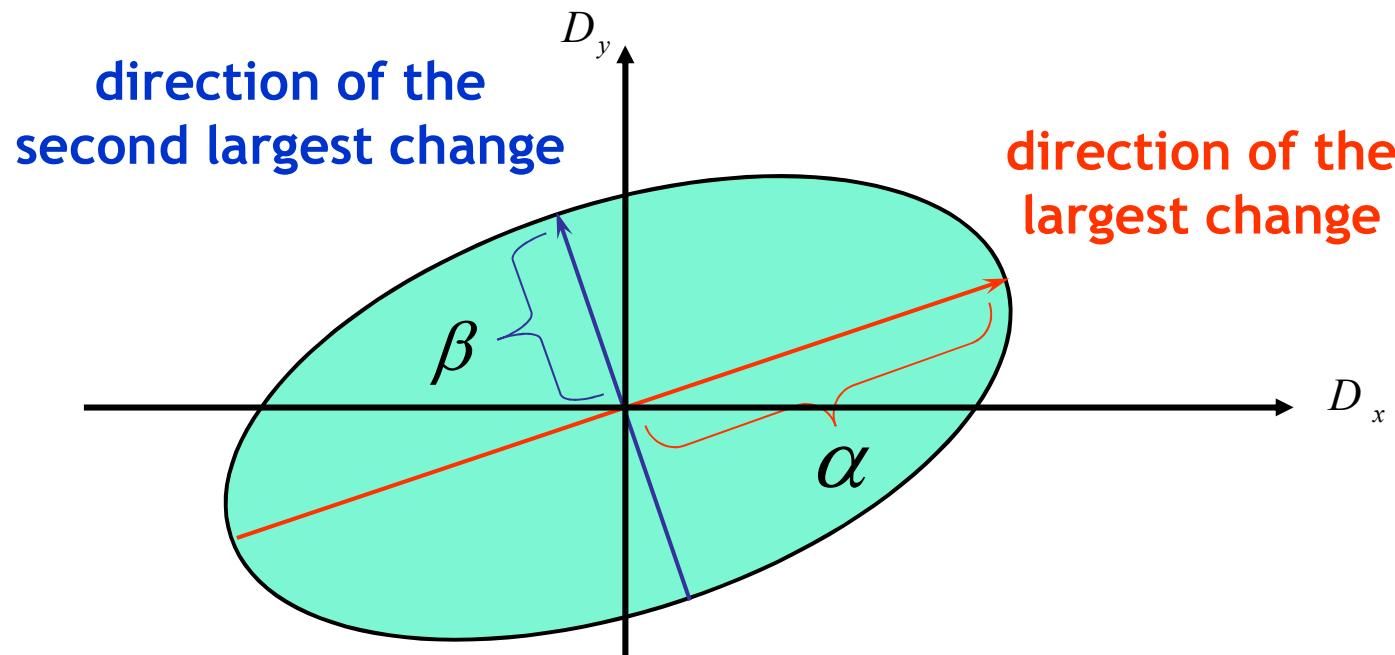
$$H = \begin{bmatrix} D_{xx} & D_{xy} \\ D_{xy} & D_{yy} \end{bmatrix} \quad \begin{aligned} \text{Tr}(H) &= D_{xx} + D_{yy} = \alpha + \beta && \text{trace} \\ \text{Det}(H) &= D_{xx}D_{yy} - (D_{xy})^2 = \alpha \beta && \text{determine} \end{aligned}$$

Hessian matrix at keypoint location

α : The larger eigenvalue of H .

β : The smaller eigenvalue of H .

$$\begin{bmatrix} D_{xx} - \alpha & D_{xy} \\ D_{xy} & D_{yy} - \beta \end{bmatrix} = 0$$



4.3 Eliminating edge responses (2/2) ??

J: Hope to keep the corner points?

$$\text{Tr}(H) = D_{xx} + D_{yy} = \alpha + \beta$$

$$\text{Det}(H) = D_{xx}D_{yy} - (D_{xy})^2 = \alpha\beta$$

$$\text{Let } \alpha = \gamma\beta$$

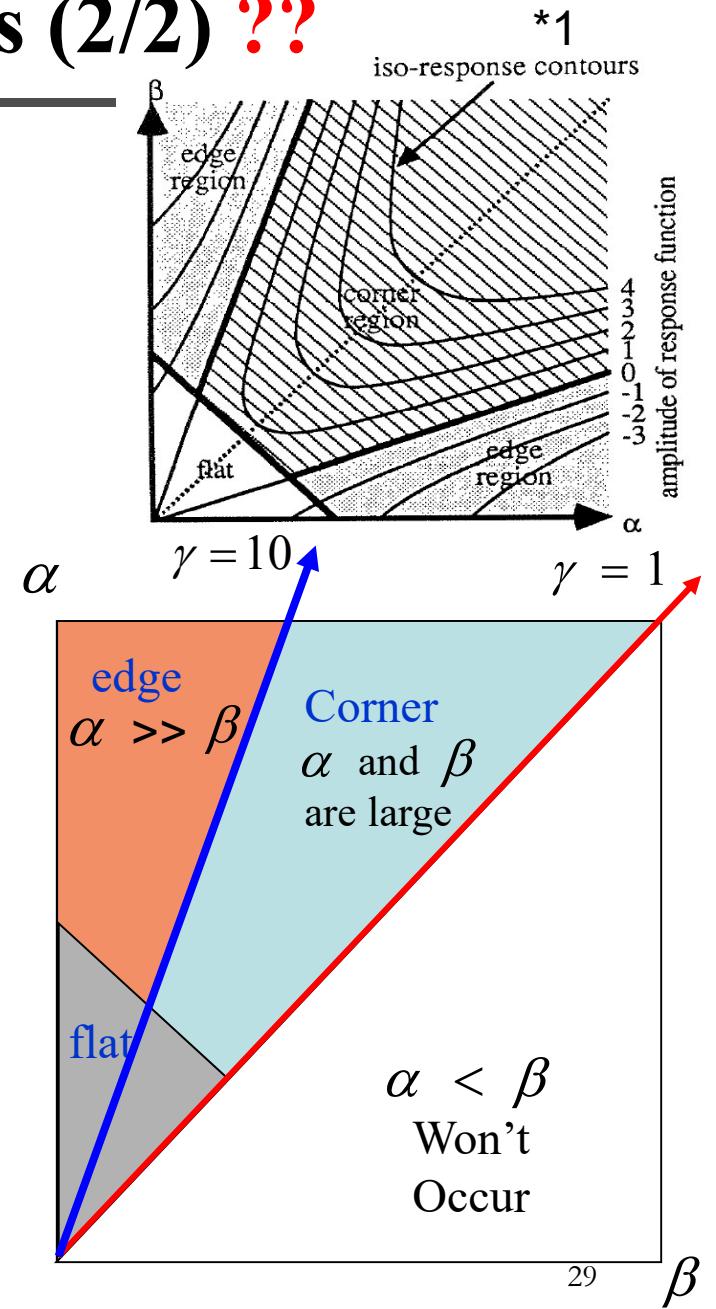
$$\frac{\text{Tr}(H)^2}{\text{Det}(H)} = \frac{(\alpha + \beta)^2}{\alpha\beta} = \frac{(\gamma\beta + \beta)^2}{\gamma\beta^2} = \frac{(\gamma + 1)^2}{\gamma}$$

Keep the (corner) points if

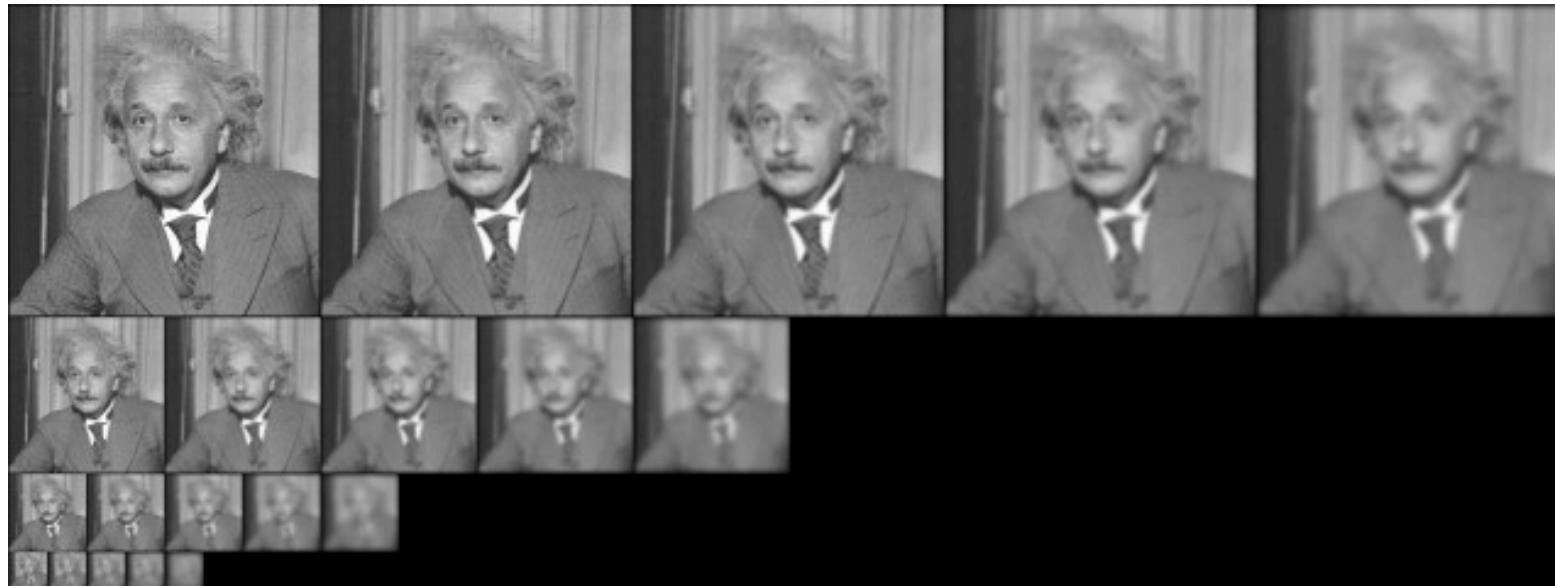
$$\gamma = 10$$

$$\frac{\text{Tr}(H)^2}{\text{Det}(H)} < \frac{(\gamma + 1)^2}{\gamma}$$

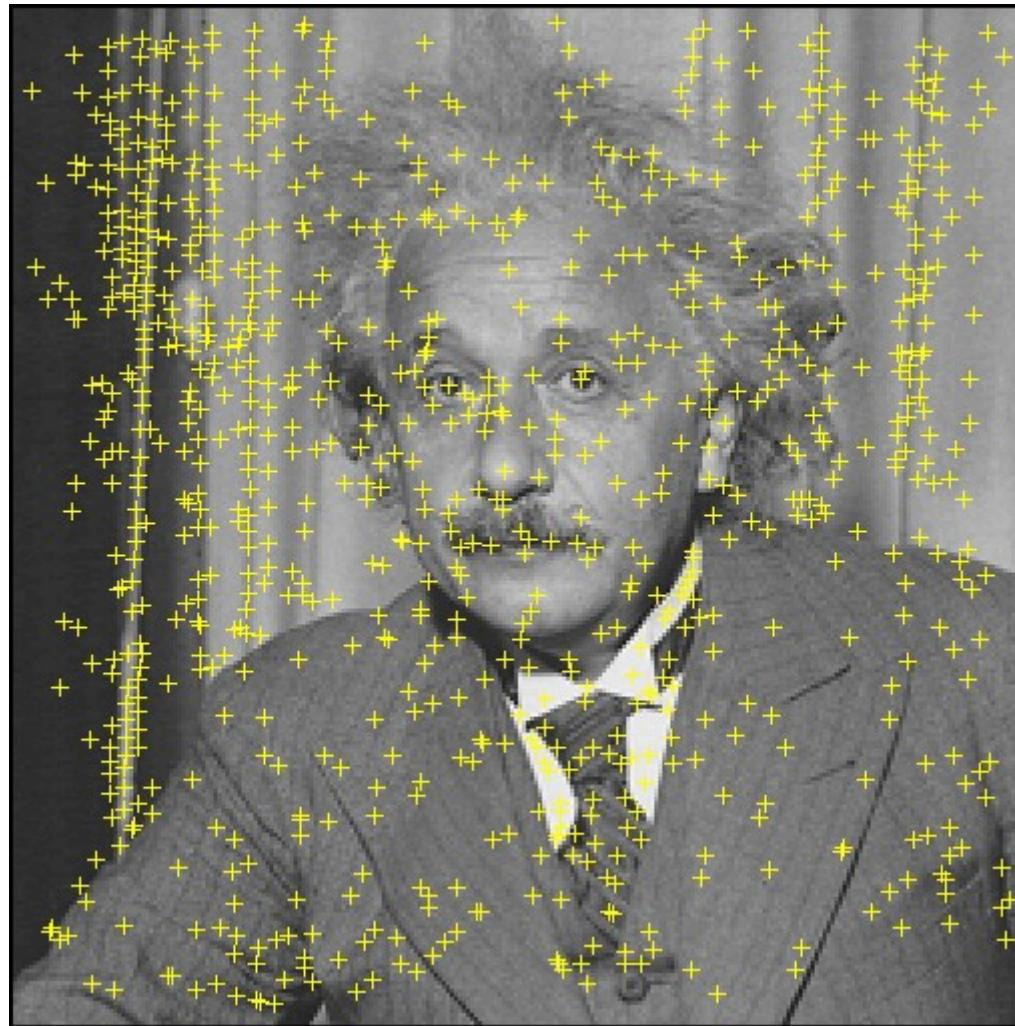
*1 C. Harris and M. Stephens, "A Combined Corner and Edge Detector,"
4th Alvey Vision Conference, pp147-151, 1988.



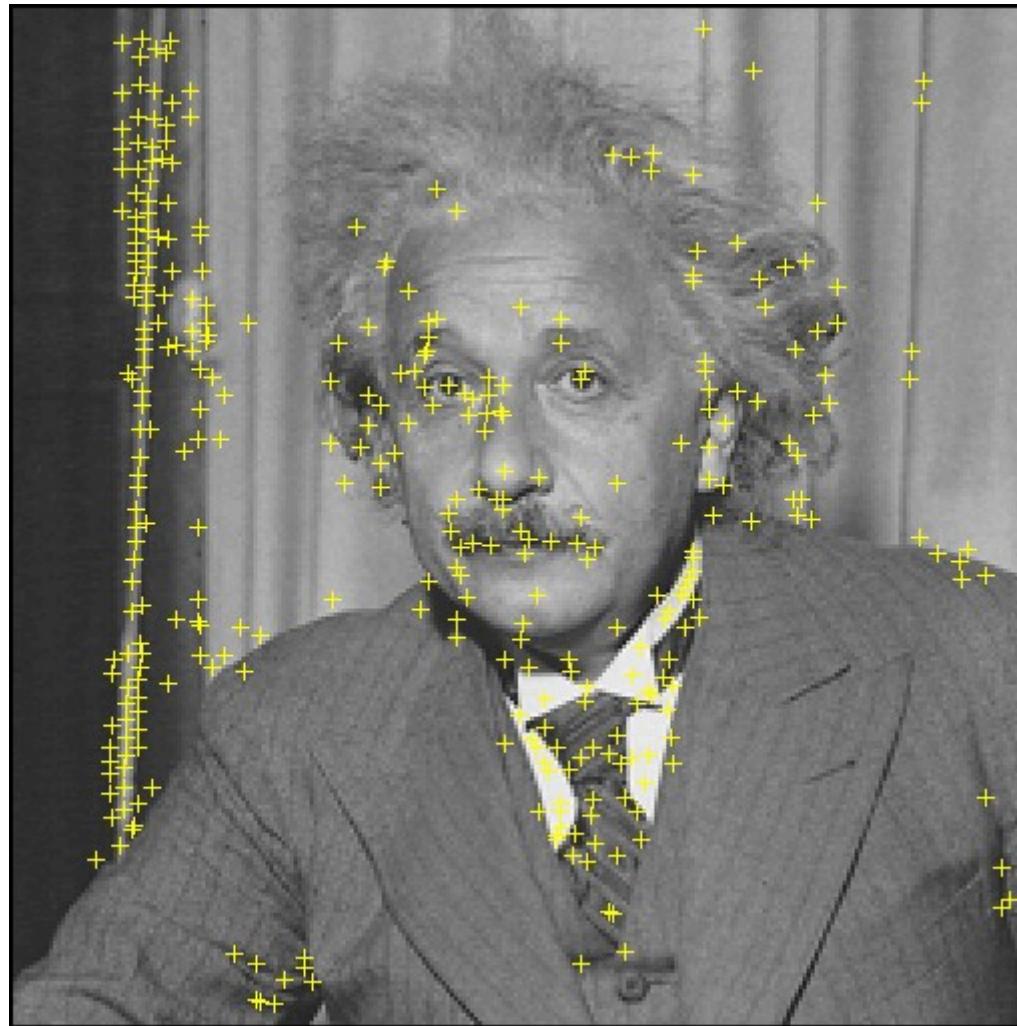
SIFT flow



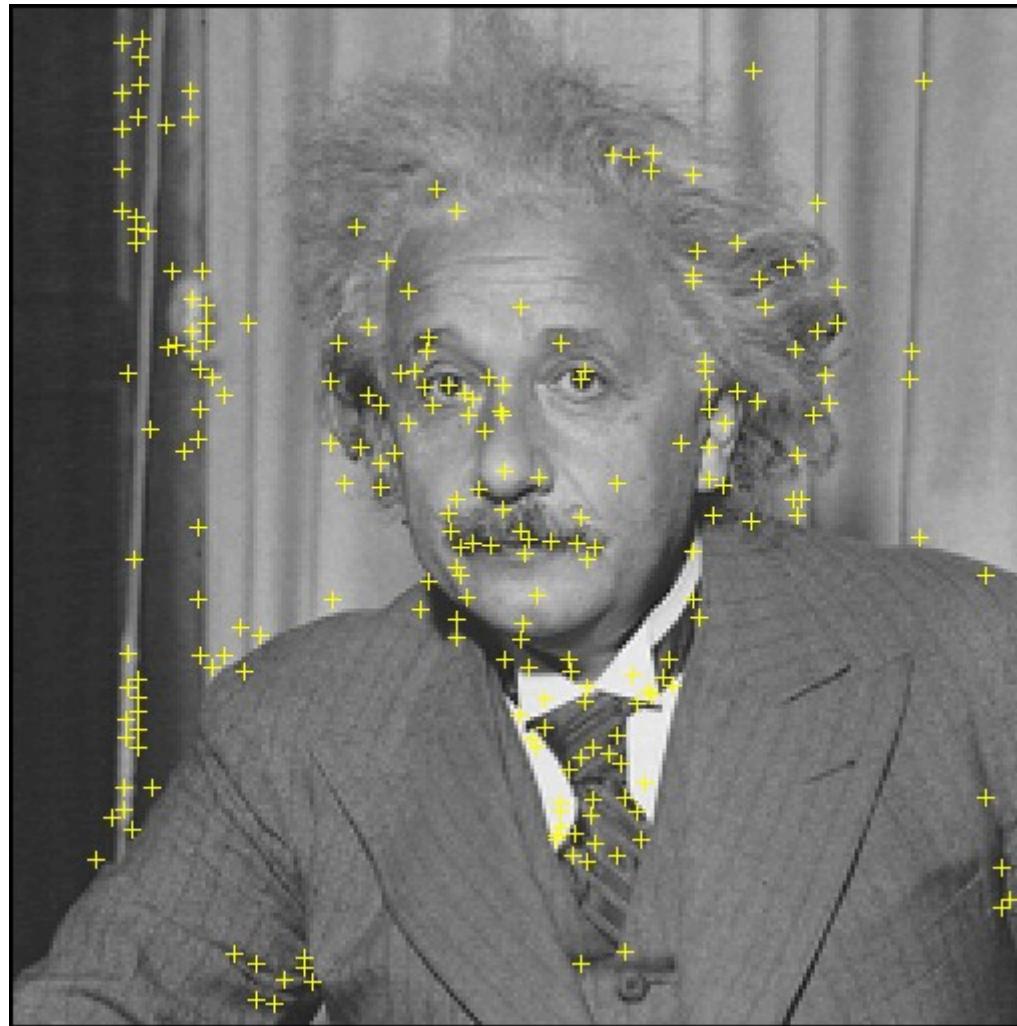
Extrema in D



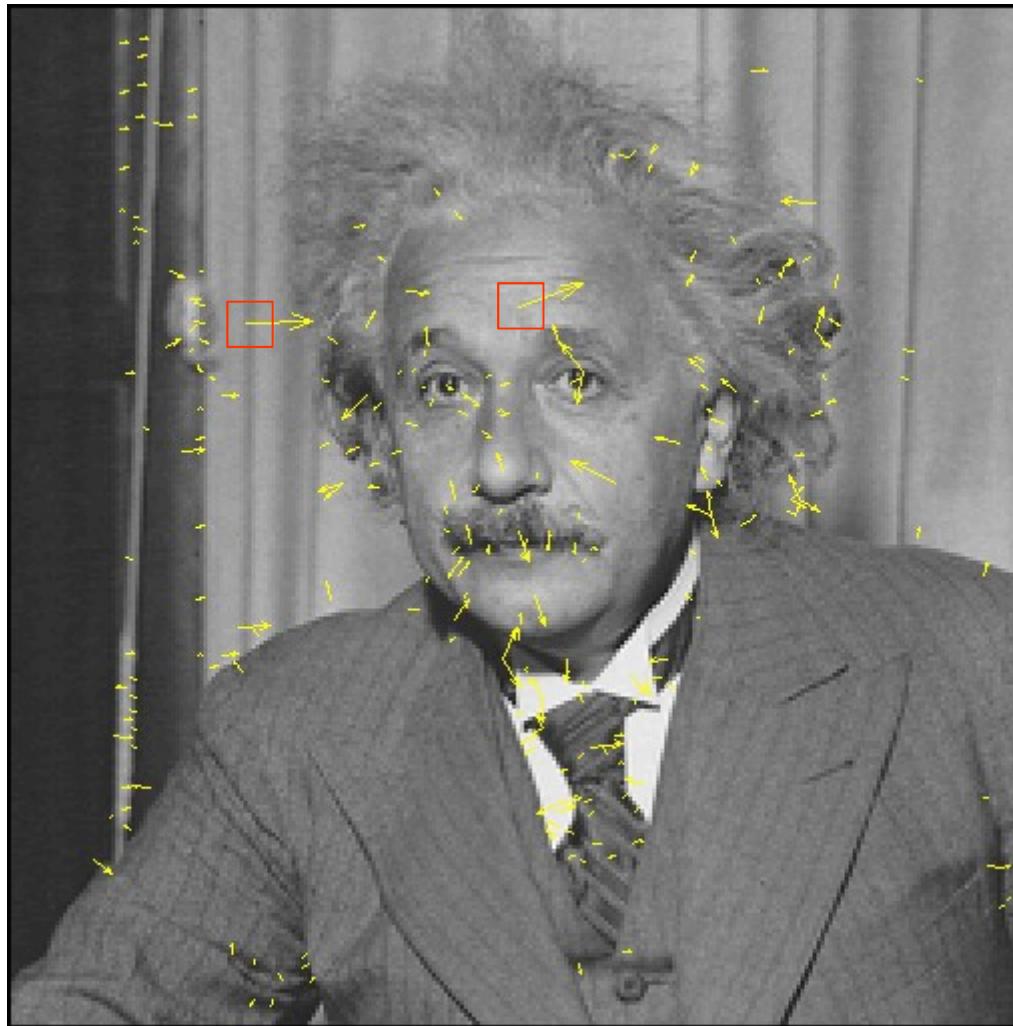
Remove low contrast



Remove edges response



SIFT descriptor





Content

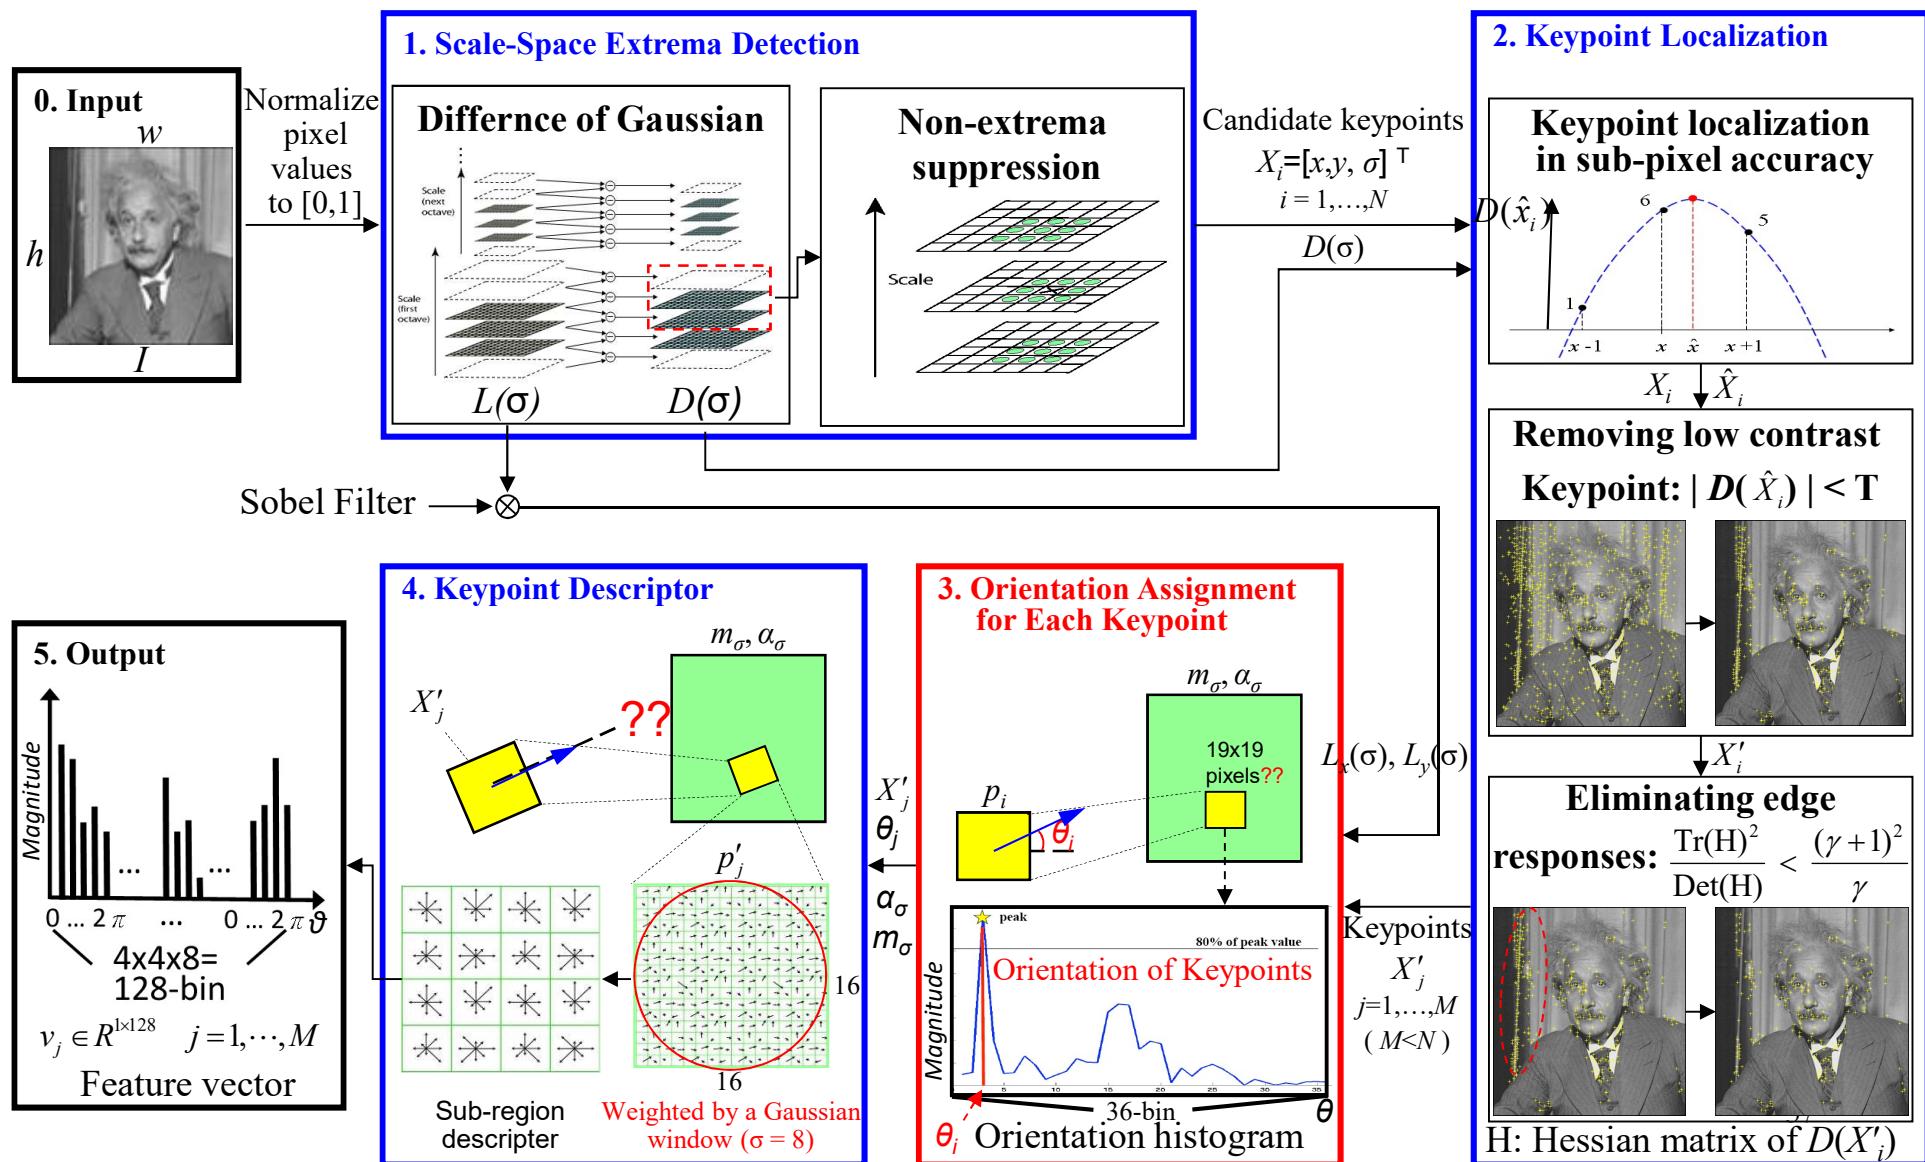
1. Motivation
2. System Flowchart
3. Detection of Scale-Space Extrema
 - 3.1 Difference of Gaussian (DoG) filter
 - 3.2 Non-maxima/minima suppression
4. Keypoint Localization
 - 4.1 Keypoint localization in sub-pixel accuracy
 - 4.2 Remove low contrast keypoints
 - 4.3 Eliminating edge responses
5. Orientation Assignment
6. Keypoint Descriptor
7. Parameter Selection

$$L(\sigma) = G(\sigma) * I$$

$$D(\sigma) = (G(k^{(n-1)}\sigma) - G(k^n\sigma)) * I$$

(Scale space)	α_σ	(Gradient orientation of $L(\sigma)$)
(DoG images)	m_σ	(Gradient magnitude of $L(\sigma)$)
	X'_i	(Accurate keypoint localization)
	θ_j	(Orientation of i -th keypoint)

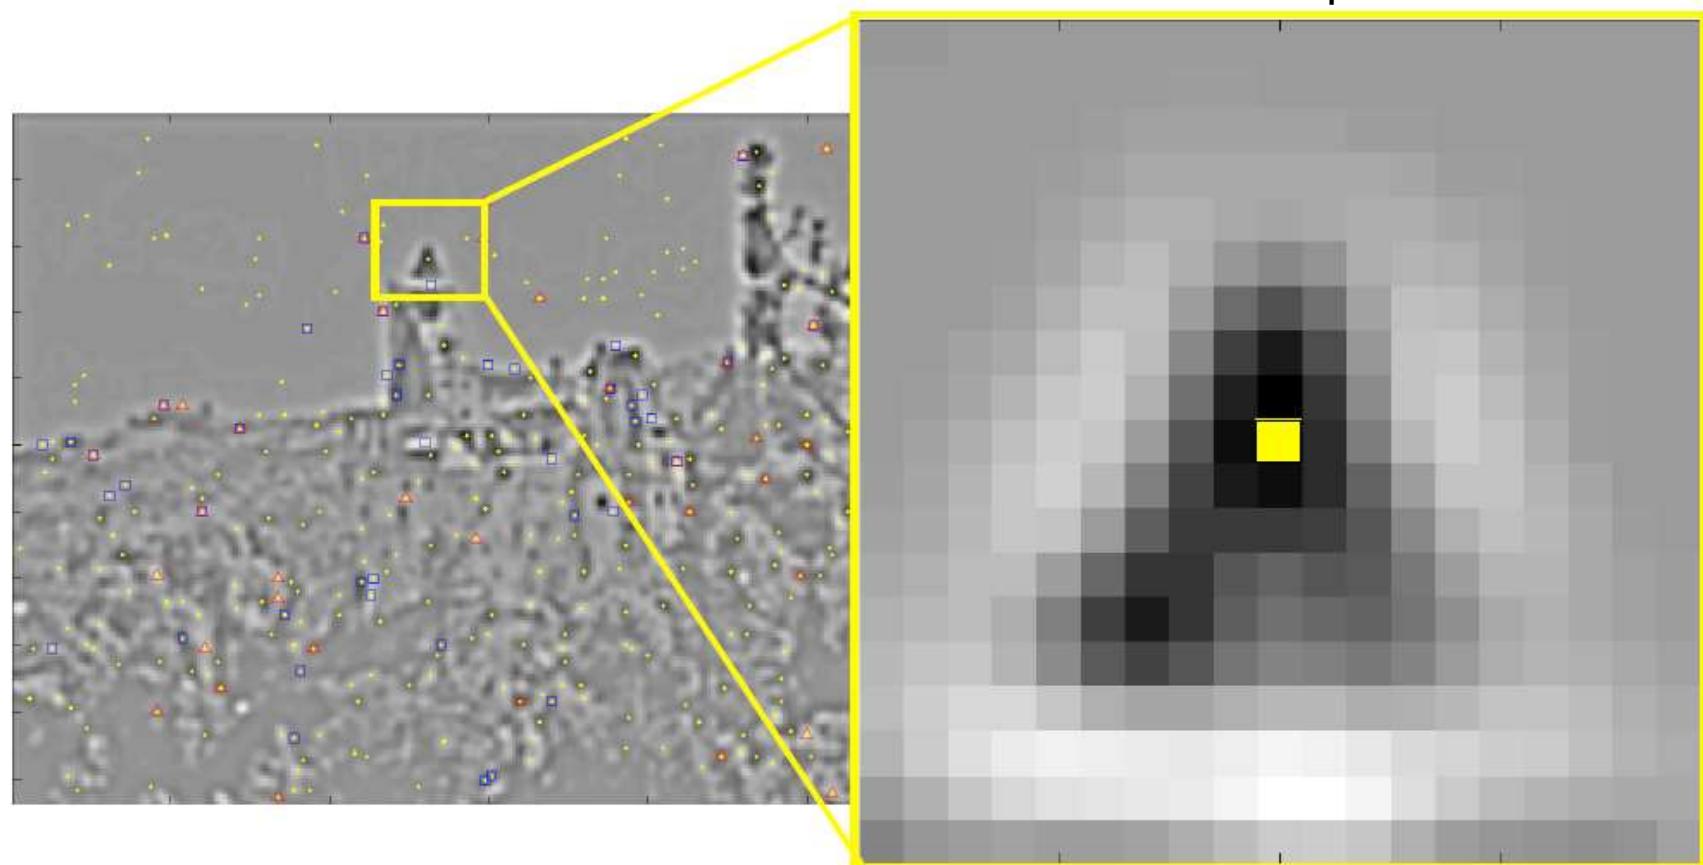
System Flowchart



5. Orientation Assignment (1/5)

- An orientation histogram is formed from the gradient orientations of sample points within a region around the keypoint X'_j (patch p_j) at $L(\sigma)$.
- By assigning a consistent orientation, the keypoint descriptor can be represented relative to this orientation and therefore achieve invariance to image rotation.
- Solve the rotation invariance problem.

5. Orientation Assignment (2/5)



- Keypoint location = extrema location
- Keypoint scale is scale of the DOG image

Previously, L is I after Gaussian smooth:

$$L(x, y, \sigma) = G(x, y, \sigma) * I(x, y)$$

Window size??,
 $\tan \Theta = 180$

5. Orientation Assignment (3/5) ??

P_j : j -th patch
 $m(x, y)$: gradient magnitude
 $\alpha(x, y)$: orientation
 P_j^m : Gradient magnitude of P_j
 P_j^α : Gradient orientation of P_j
 $L(\sigma)$: Scale space of image.

- Each sample added to the histogram is weighted by its **gradient magnitude** and by a **Gaussian-weighted circular window**.

$$L_x(x, y, \sigma) = L(x+1, y, \sigma) - L(x-1, y, \sigma) \Rightarrow m_\sigma(x, y) = \sqrt{L_x(x, y, \sigma)^2 + L_y(x, y, \sigma)^2}$$

$$L_y(x, y, \sigma) = L(x, y+1, \sigma) - L(x, y-1, \sigma)$$

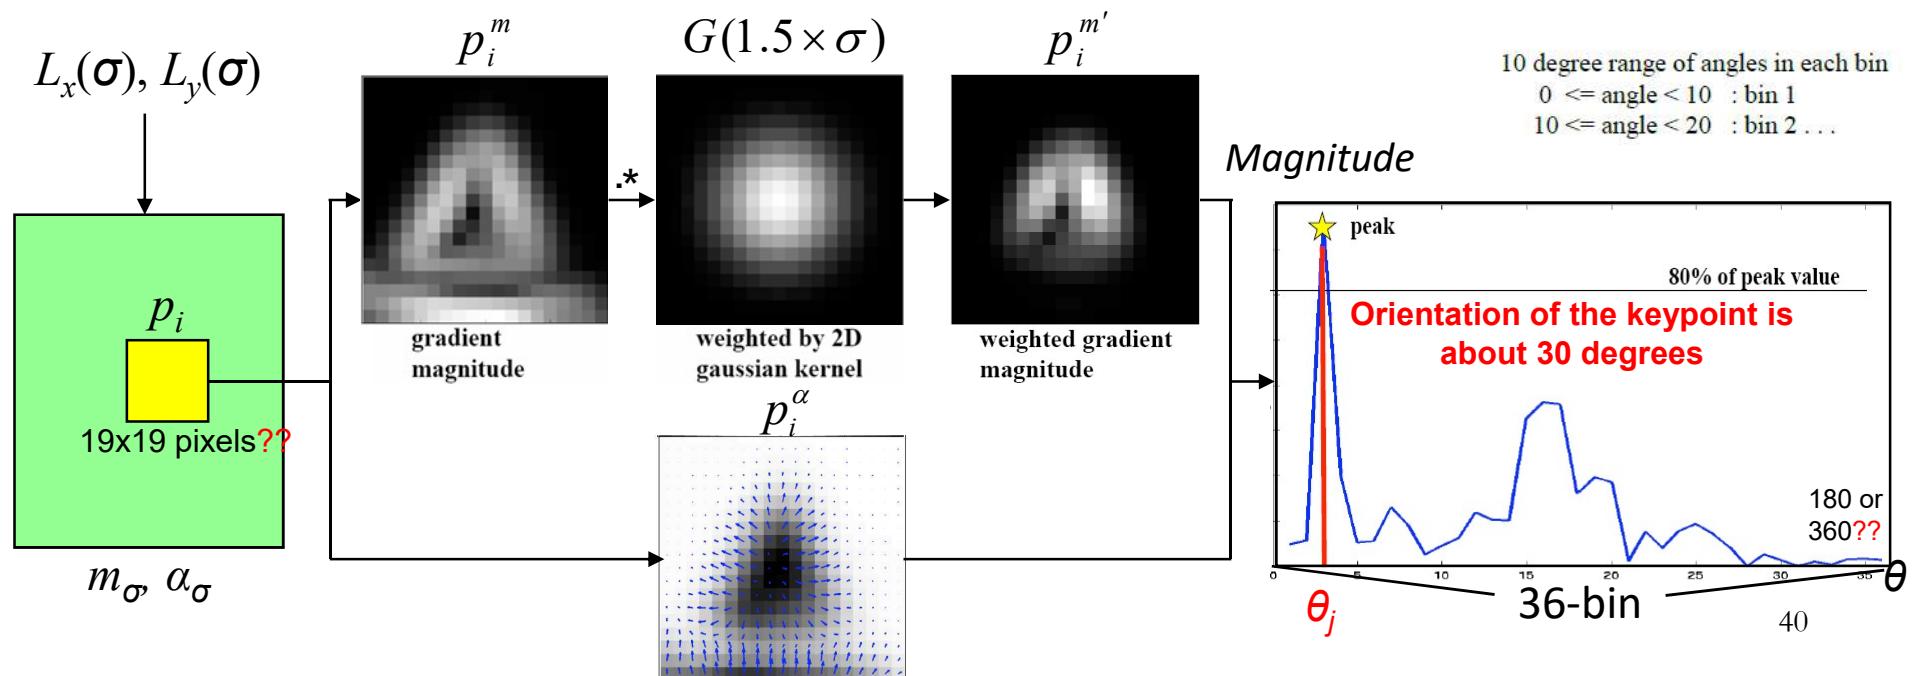
$$L_x(\sigma) = \begin{bmatrix} -1 & 0 & +1 \\ x-1 & x & x+1 \end{bmatrix} * L$$

$$L_y(\sigma) = \begin{bmatrix} -1 & y-1 \\ 0 & y \\ +1 & y+1 \end{bmatrix} * L$$

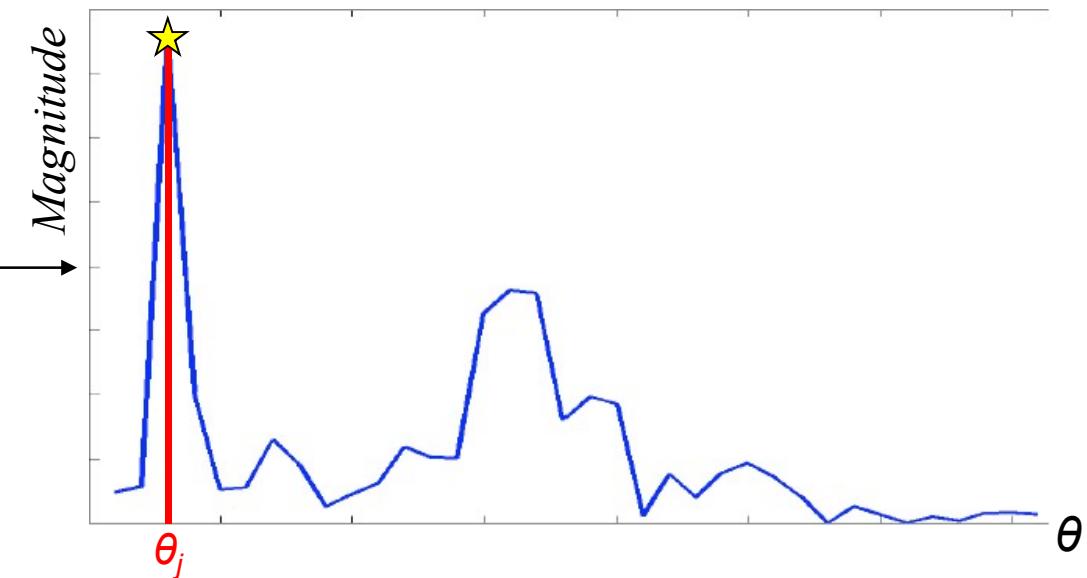
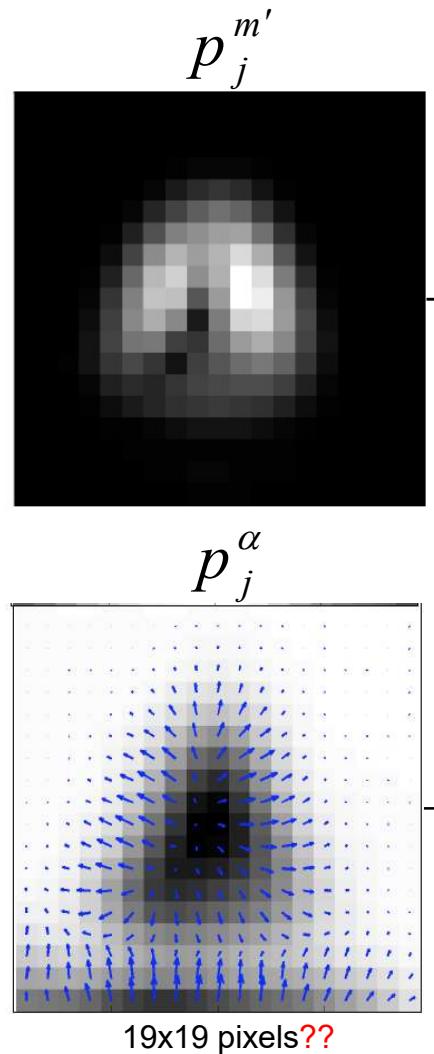
➢ Consistent orientation for each feature is θ_j .

$$\alpha_\sigma(x, y) = \tan^{-1}\left(\frac{L_y(x, y, \sigma)}{L_x(x, y, \sigma)}\right)$$

$$p_j^{m'}(x, y) = p_j^m(x, y) \times G(x, y, 1.5 \times \sigma)$$



5. Orientation Assignment (4/5)



36 bins

10 degree range of angles in each bin

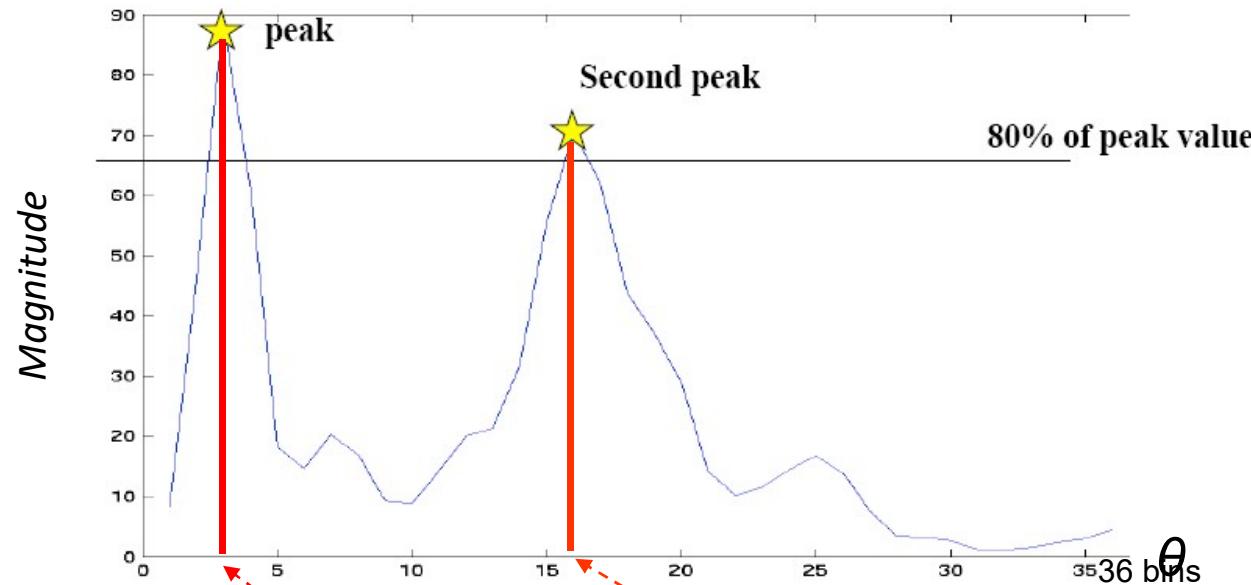
$0 \leqslant \text{angle} < 10$: bin 1

$10 \leqslant \text{angle} < 20$: bin 2

$20 \leqslant \text{angle} < 30$: bin 3 ...

5. Orientation Assignment (5/5)

There may be multiple orientations.



In this case, generate duplicate keypoints, one with orientation at 30 degrees, one at 155 degrees.

- Any other local peak that is **within** 80% of the highest peak is used to also create a keypoint with that orientation.

Content

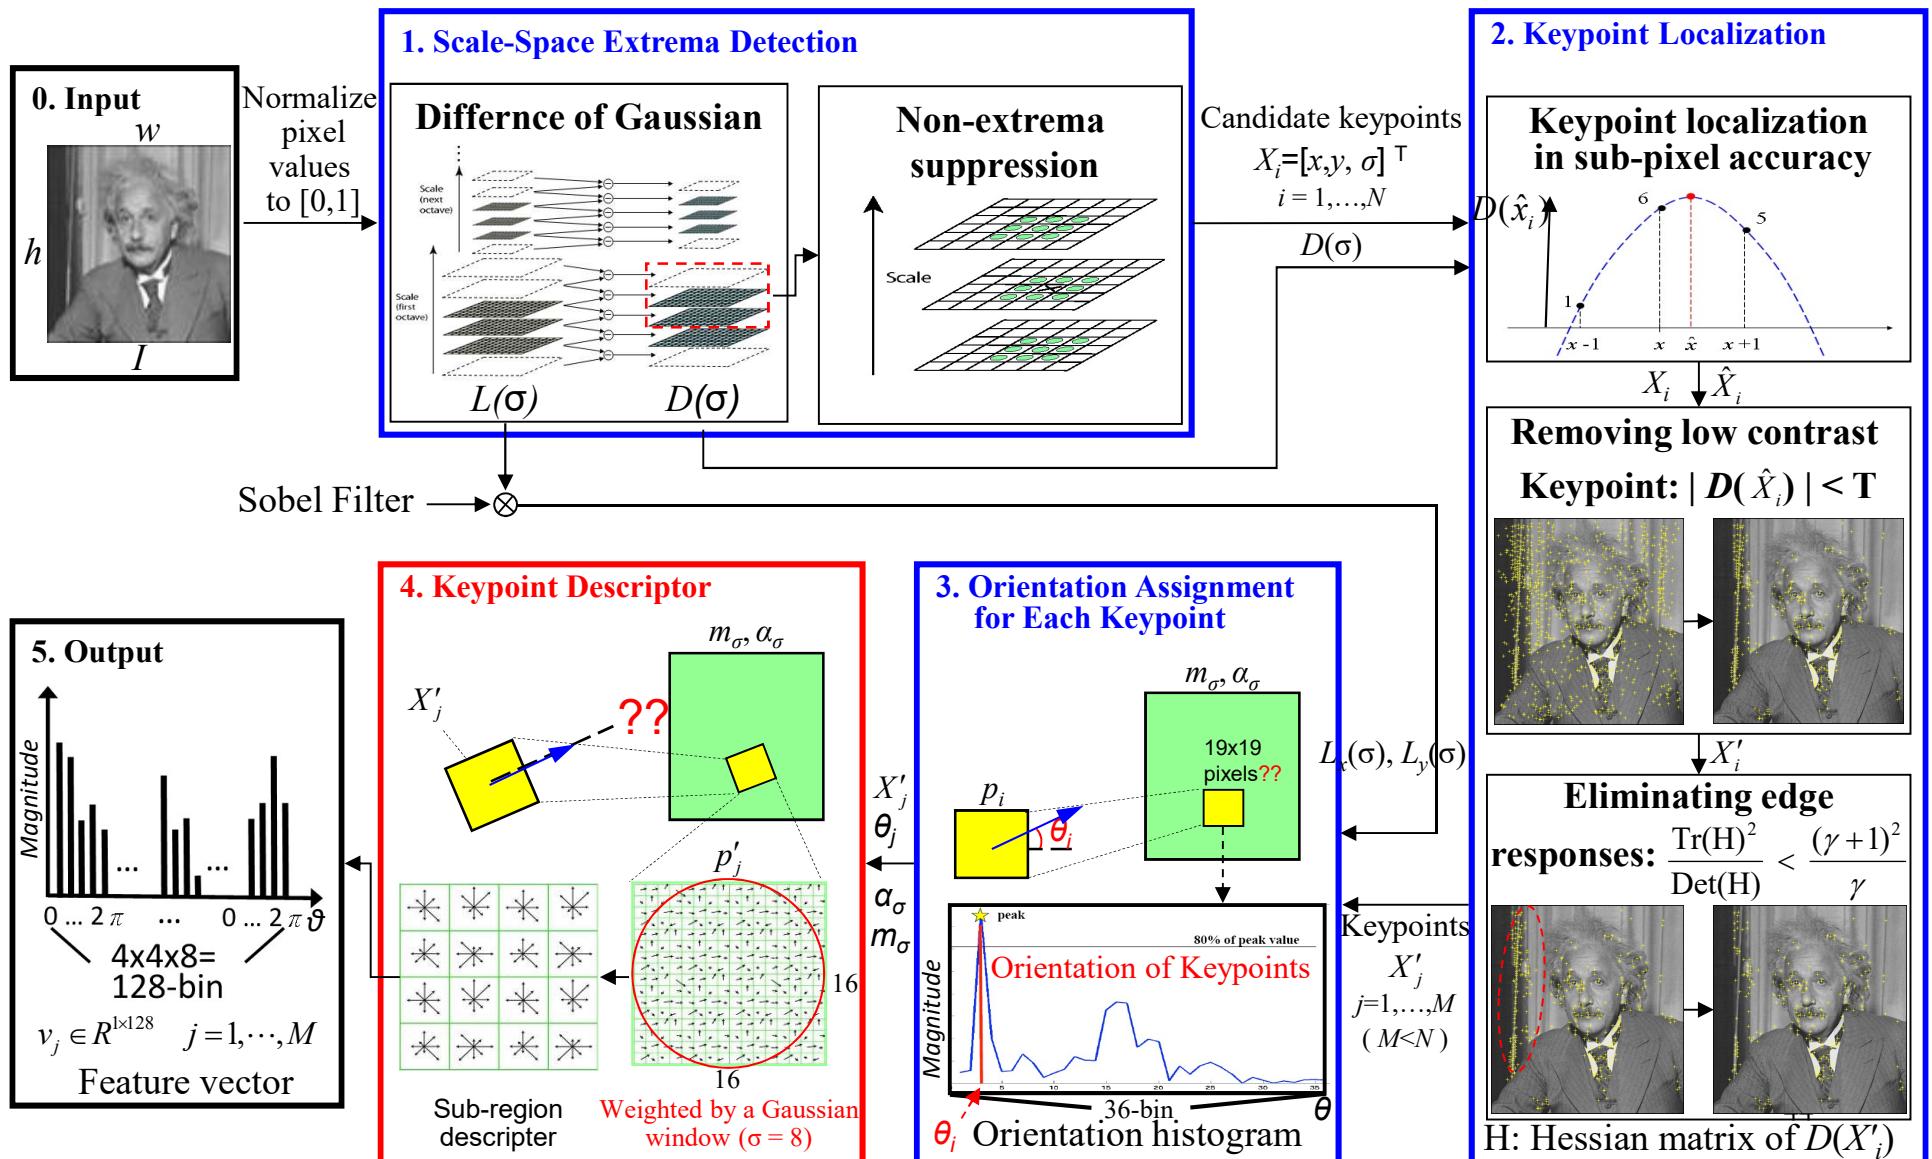
1. Motivation
2. System Flowchart
3. Detection of Scale-Space Extrema
 - 3.1 Difference of Gaussian (DoG) filter
 - 3.2 Non-maxima/minima suppression
4. Keypoint Localization
 - 4.1 Keypoint localization in sub-pixel accuracy
 - 4.2 Remove low contrast keypoints
 - 4.3 Eliminating edge responses
5. Orientation Assignment
6. Keypoint Descriptor
7. Parameter Selection

$$L(\sigma) = G(\sigma) * I$$

$$D(\sigma) = (G(k^{(n-1)}\sigma) - G(k^n\sigma)) * I$$

(Scale space)	α_σ	(Gradient orientation of $L(\sigma)$)
(DoG images)	m_σ	(Gradient magnitude of $L(\sigma)$)
	X'_i	(Accurate keypoint localization)
	θ_j	(Orientation of i -th keypoint)

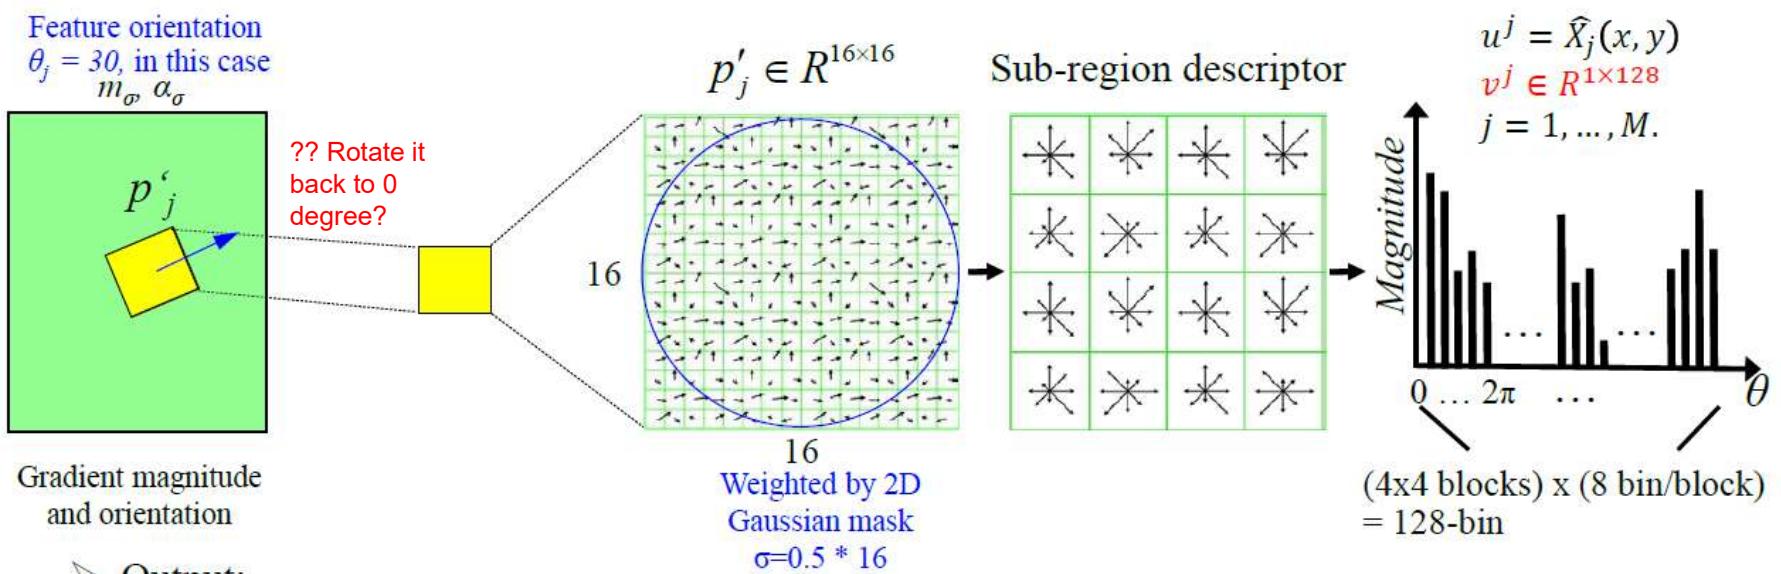
System Flowchart



6. Keypoint Descriptor ??

4) Compute feature descriptor.

- Image gradients are sampled over 16x16 array of locations **in scale space** relative to the **keypoint orientation**.
- The feature vector, v_j (*128 dimensions*), is **normalized to unit length**. (Reduce change contrast) $v_j \rightarrow \frac{v_j}{\|v_j\|}$



- Output:
 - Each image has M features.
 - Each feature has 128 dimension feature vector.
 - The center position of each feature.

$$u_n^j = [x, y]$$

$$v_n^j \in R^{1 \times 128}$$

$$k = 1, \dots, M.$$

: The center position of feature j in image n.

: Feature vector j in image n (Feature Descriptor).

Content

1. Motivation

2. System Flowchart

3. Detection of Scale-Space Extrema

 3.1 Difference of Gaussian (DoG) filter

 3.2 Non-maxima/minima suppression

4. Keypoint Localization

 4.1 Keypoint localization in sub-pixel accuracy

 4.2 Remove low contrast keypoints

 4.3 Eliminating edge responses

5. Orientation Assignment

6. Keypoint Descriptor

7. Parameter Selection

7. Parameter Selection (1/2)

- We use a collection of **32 real images** drawn from a diverse range, including (image domain) outdoor scenes, human faces, aerial photographs, and industrial images.
- The **image domain** was found to have almost no influence on any of the results.
- Each image was then subject to **a range of transformations**, including **rotation** (random), **scaling** (0.2 ~ 0.9), **change in brightness** and **contrast**, and **addition of image noise** (1%).

7. Parameter Selection (2/2)

- We define a **matching scale** as being within a factor of $\sqrt{2}$ of the correct scale, and a **matching location** as being within σ pixels, where σ is the scale of the keypoint.
- Distinguishable:
 - For a feature x , we found the **closest feature x_1** and the **second closest feature x_2** .
 - If the distance **ratio** of $d(x, x_1)$ and $d(x, x_2)$ is smaller than **0.8**, then it is accepted as a match (x and x_1 are the corresponding point pair).

Merge to
one??

7. Parameter Selection: s

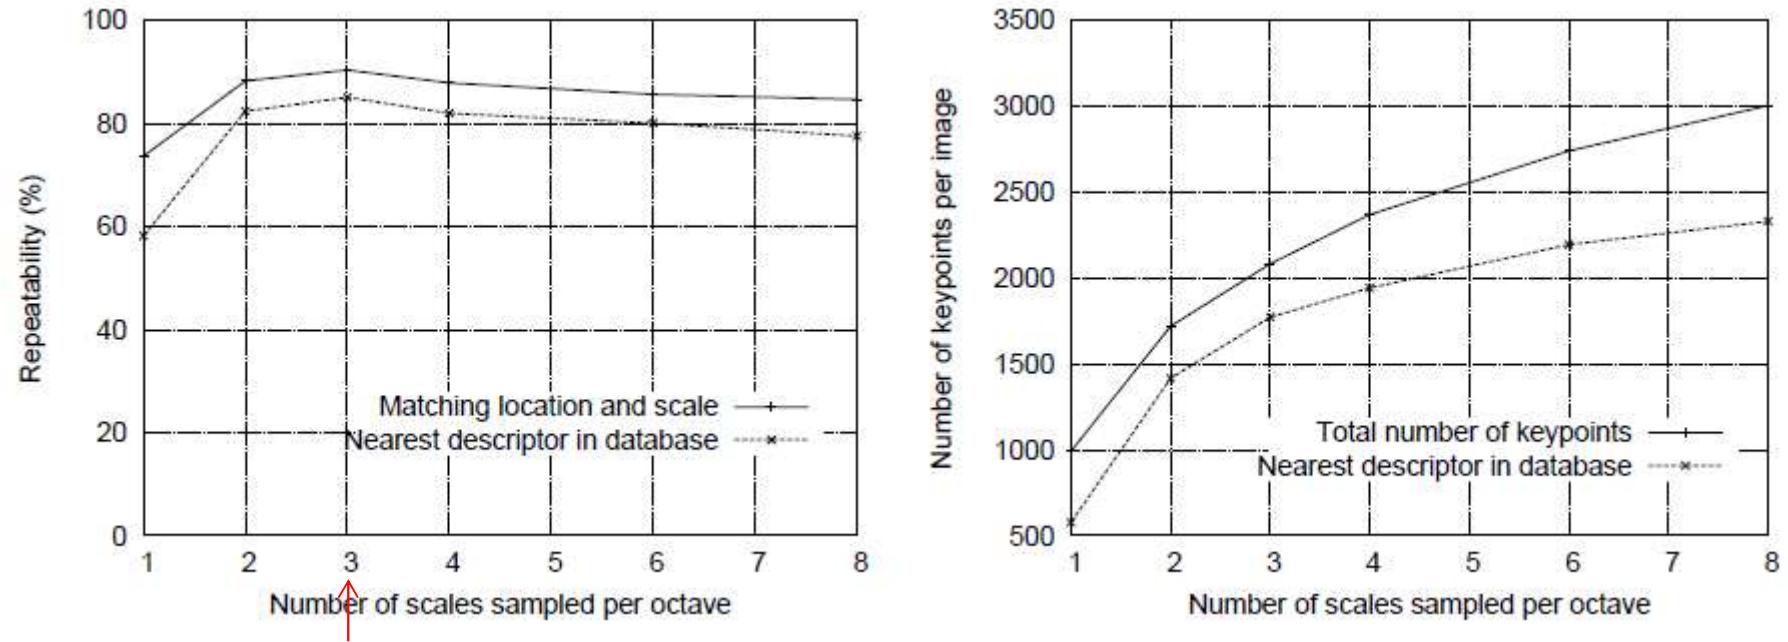


Figure 3:

The top line of the first graph shows the percent of keypoints that are repeatably detected at the same location and scale in a transformed image as a function of the number of scales sampled per octave. The lower line shows the percent of keypoints that have their descriptors correctly matched to a large database. The second graph shows the total number of keypoints detected in a typical image as a function of the number of scale samples.

7. Parameter Selection: Pre-smoothing

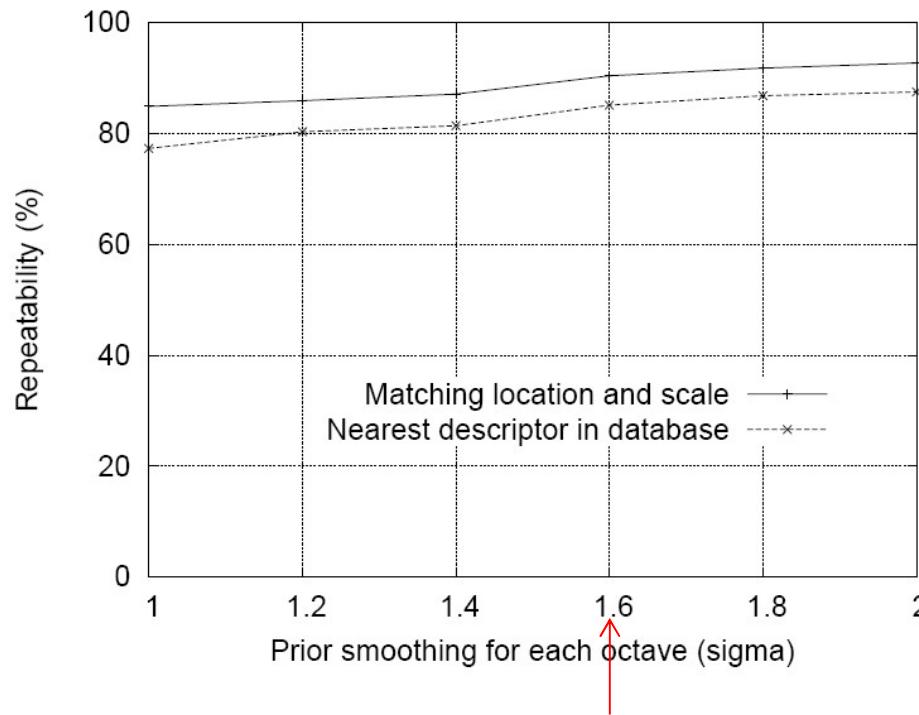


Figure 4: The top line in the graph shows the percent of keypoint locations that are repeatably detected in a transformed image as a function of the prior image smoothing for the first level of each octave. The lower line shows the percent of descriptors correctly matched against a large database.

There is a cost to using a large σ in terms of efficiency.

We have chosen to use $\sigma = 1.6$, which provides close to optimal repeatability

7. Parameter Selection: Why 4x4x8?

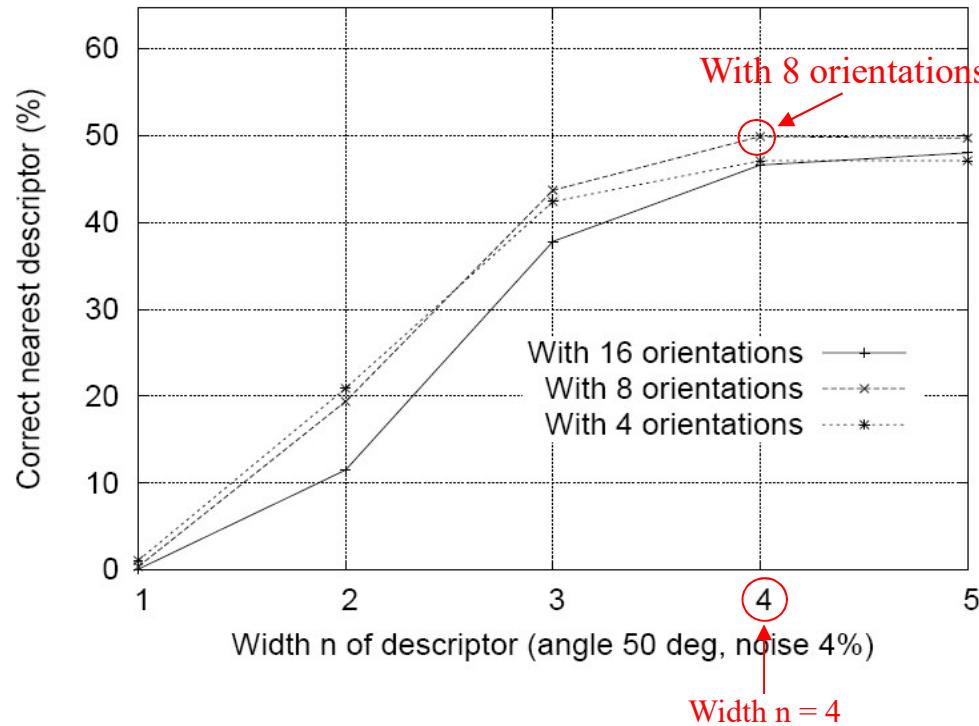


Figure 8: This graph shows the percent of keypoints giving the correct match to a database of 40,000 keypoints as a function of width of the $n \times n$ keypoint descriptor and the number of orientations in each histogram. The graph is computed for images with affine viewpoint change of 50 degrees and addition of 4% noise.

The graph shows that a single orientation histogram ($n = 1$) is very poor at discriminating, but the results continue to improve up to a **4x4 array** of histograms with **8 orientations**.

7. Sensitivity to affine change

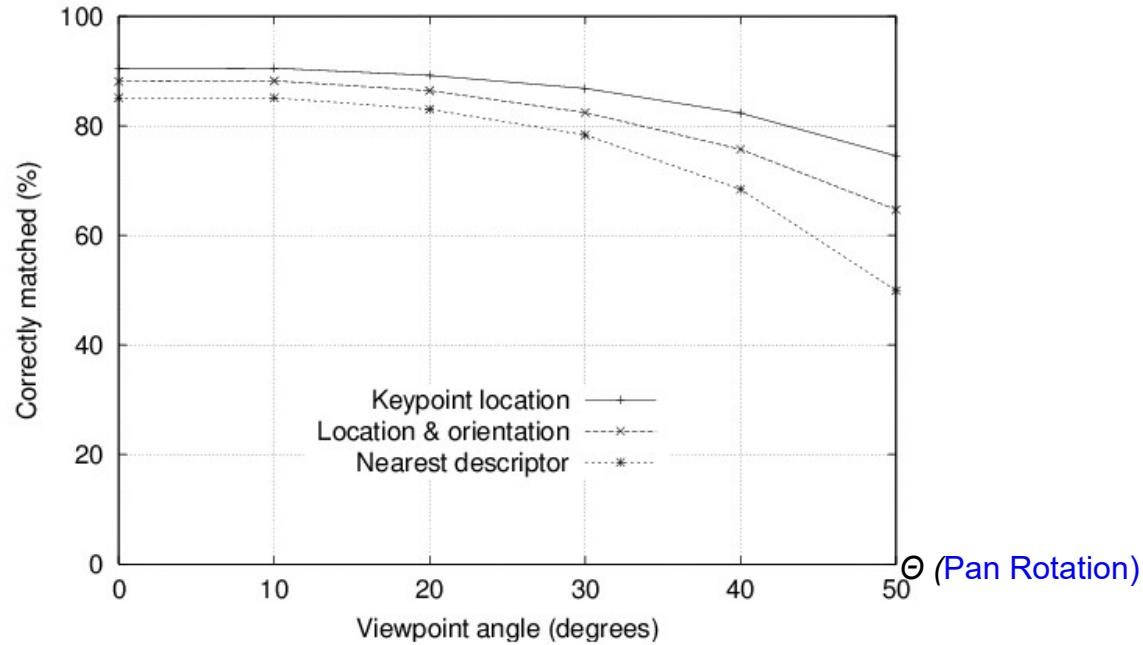
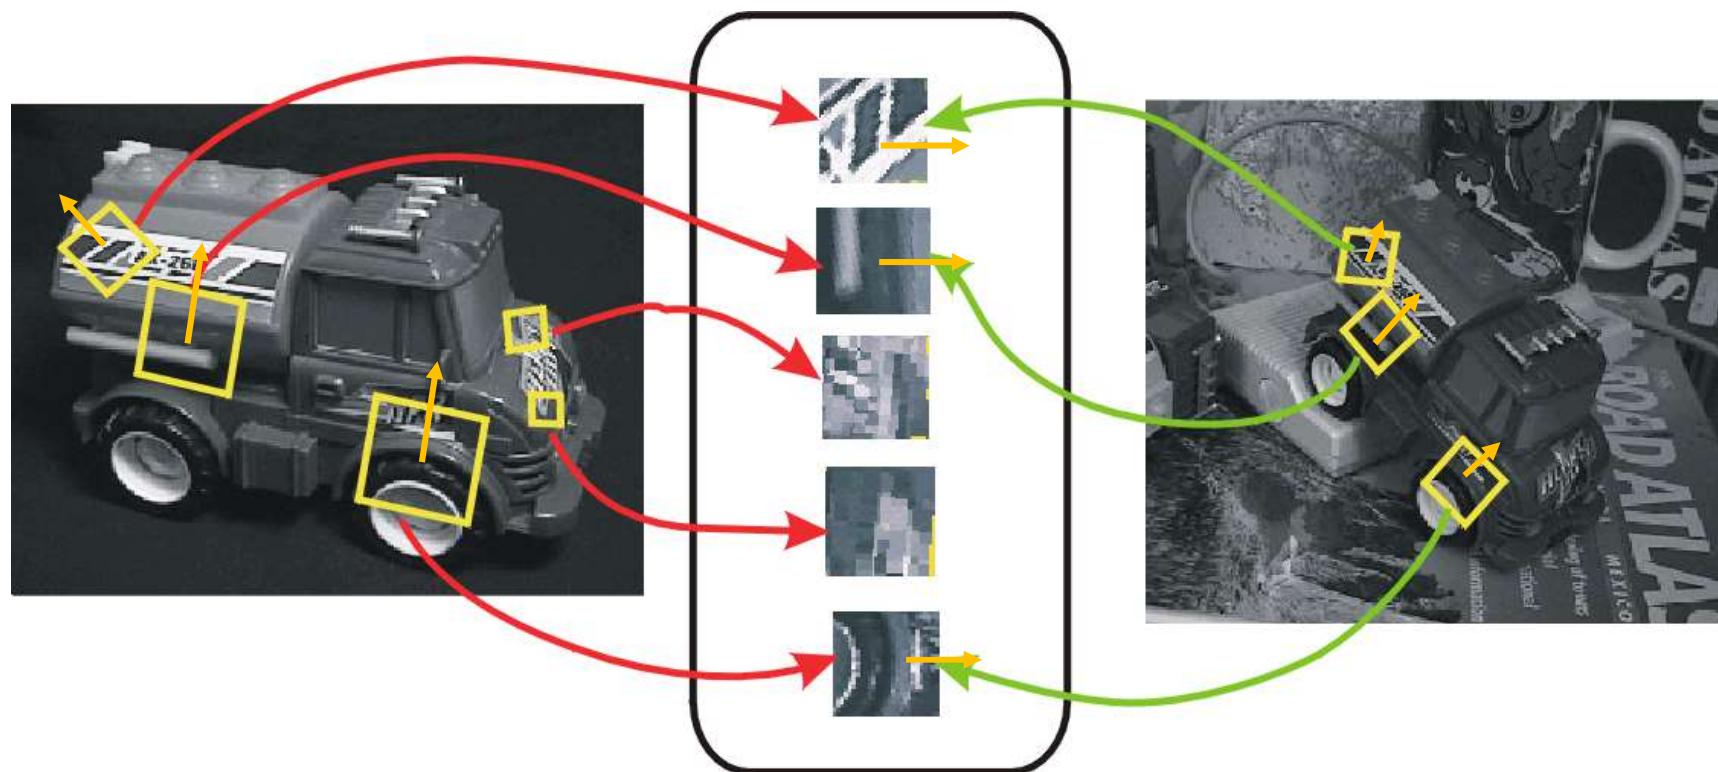
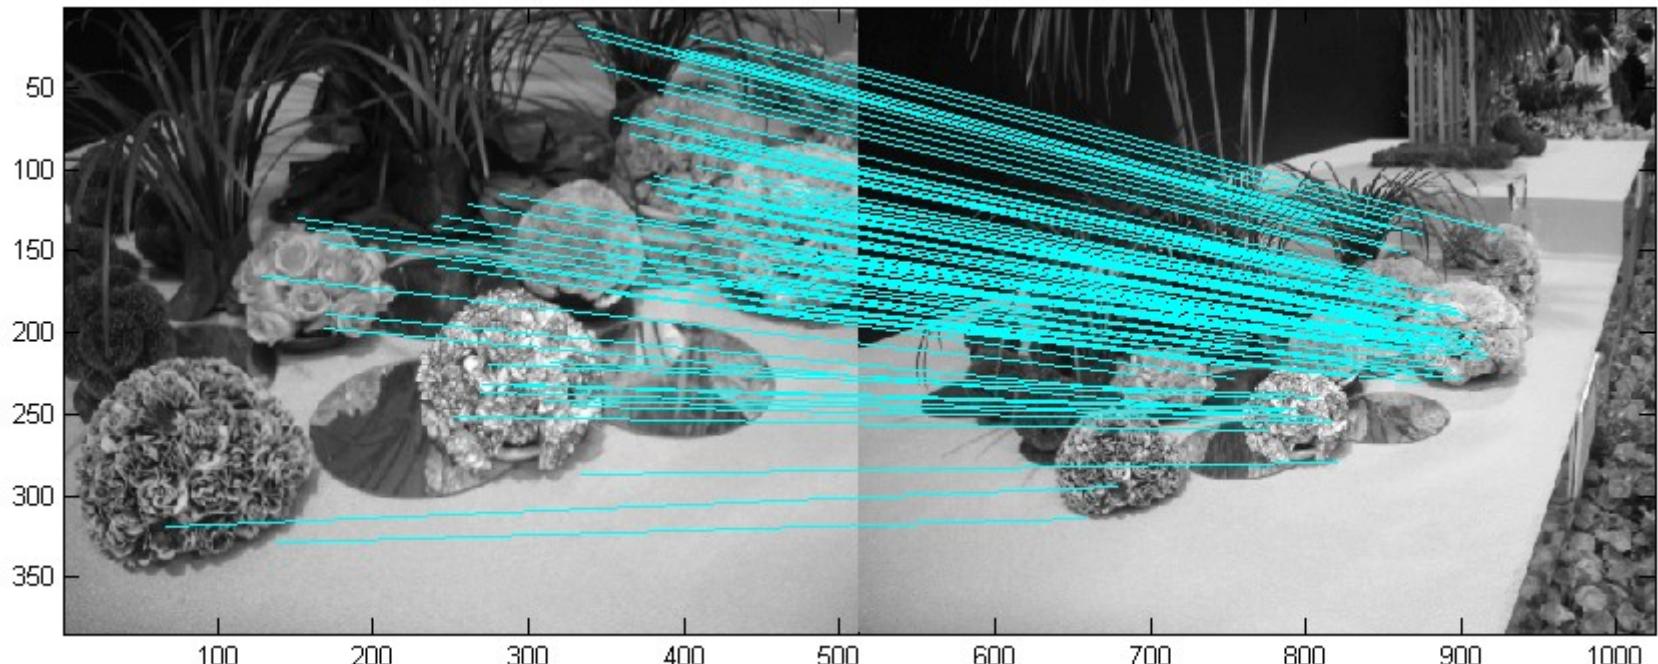


Figure 9: This graph shows the stability of detection for keypoint location, orientation, and final matching to a database as a function of affine distortion. The degree of affine distortion is expressed in terms of the equivalent viewpoint rotation in depth for a planar surface.

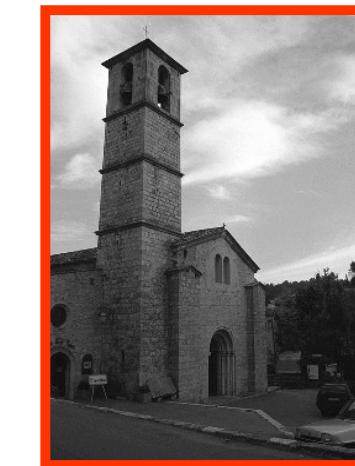
8.1 Example



8.1 Example



8.2 Image retrieval

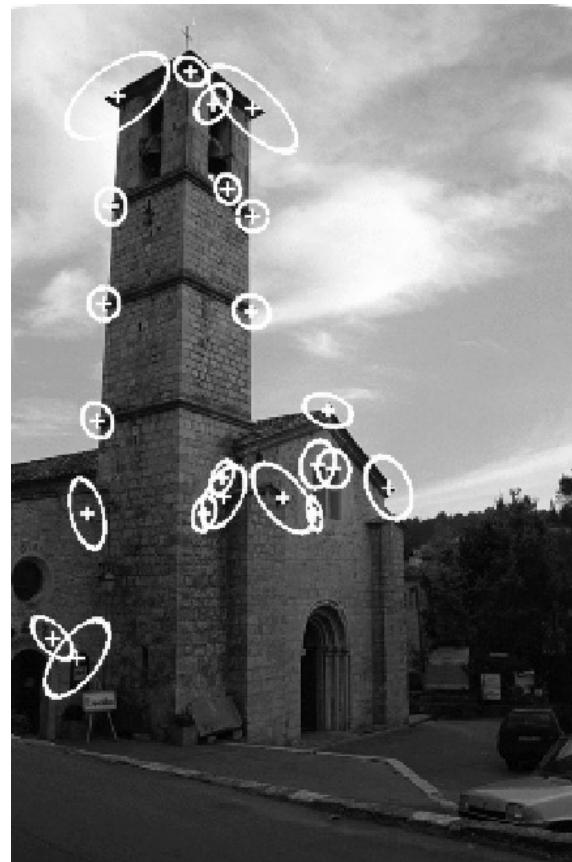
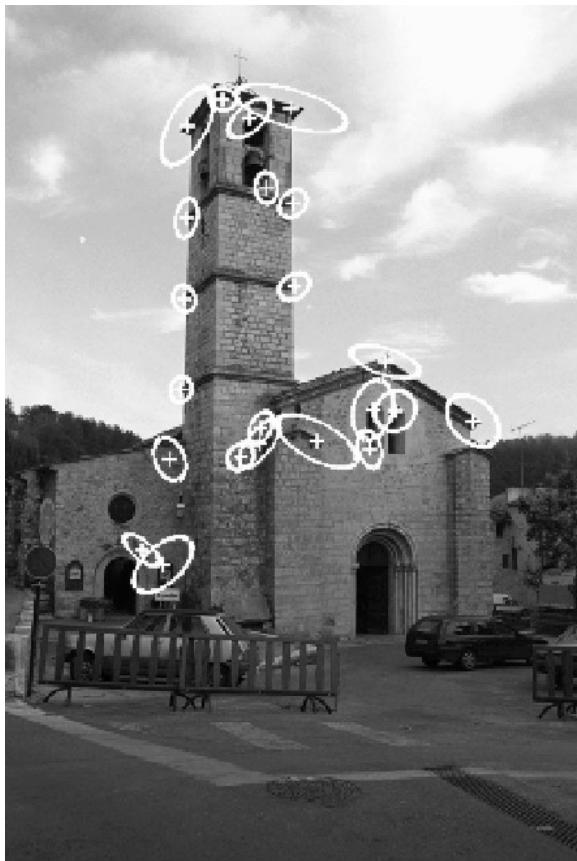


change in viewing angle



• • •
> 5000
images

8.2 Image retrieval



22 correct matches

8.2 Image retrieval



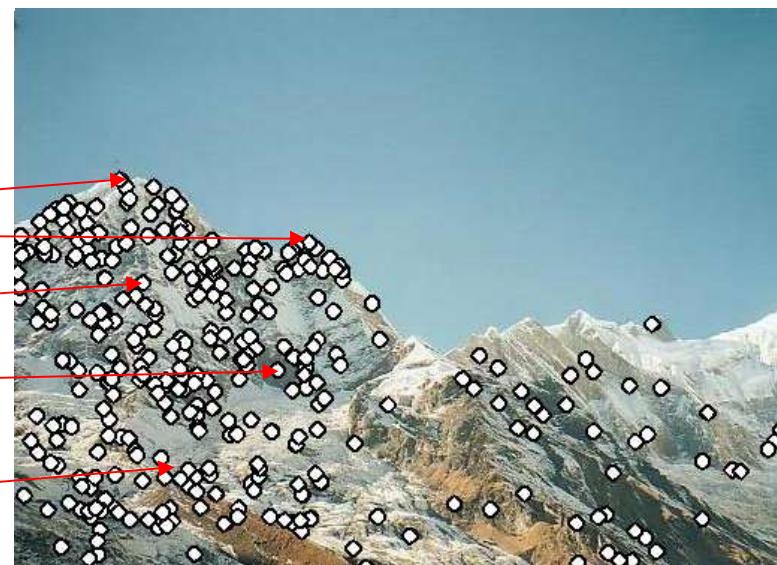
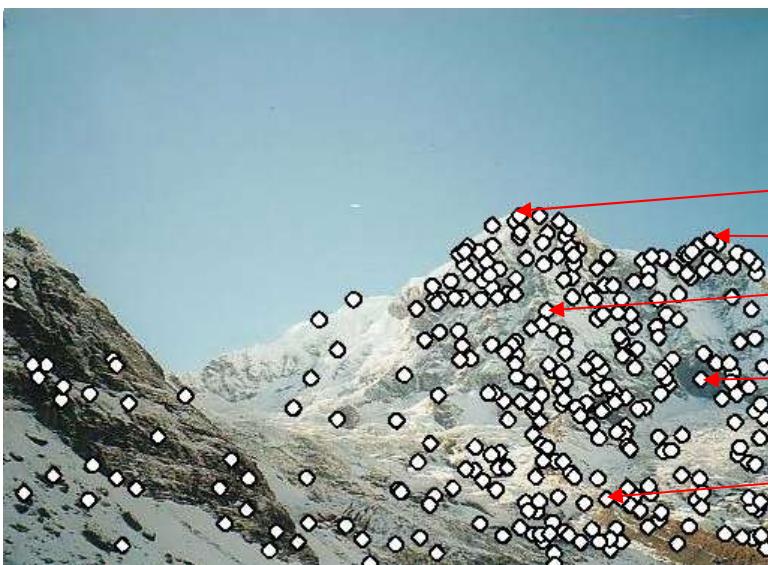
change in viewing angle
+ scale change



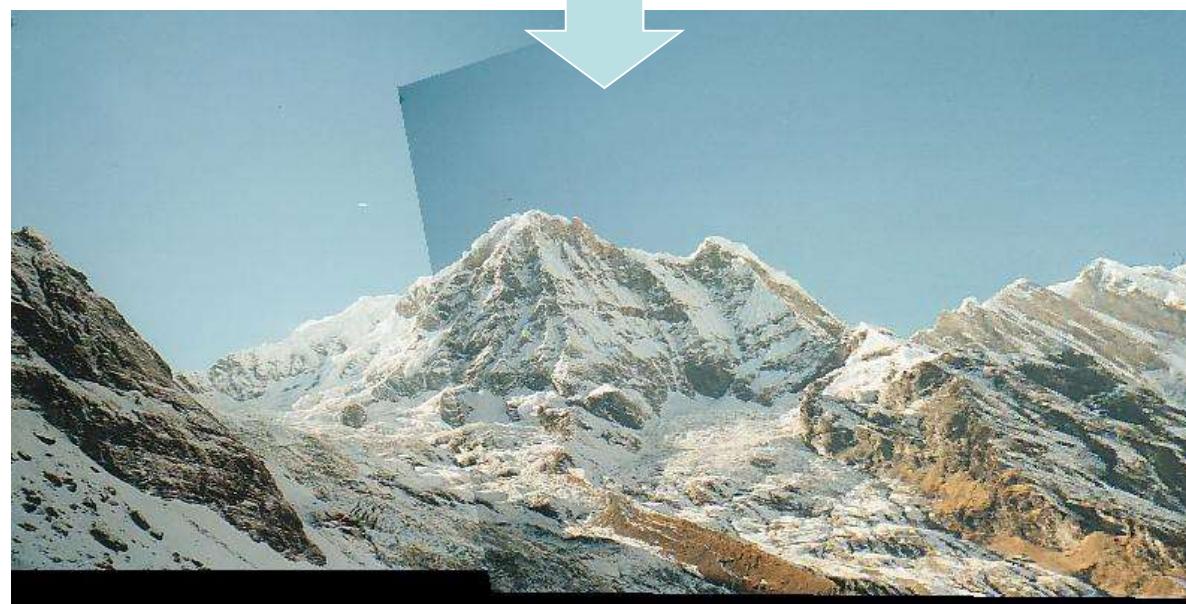
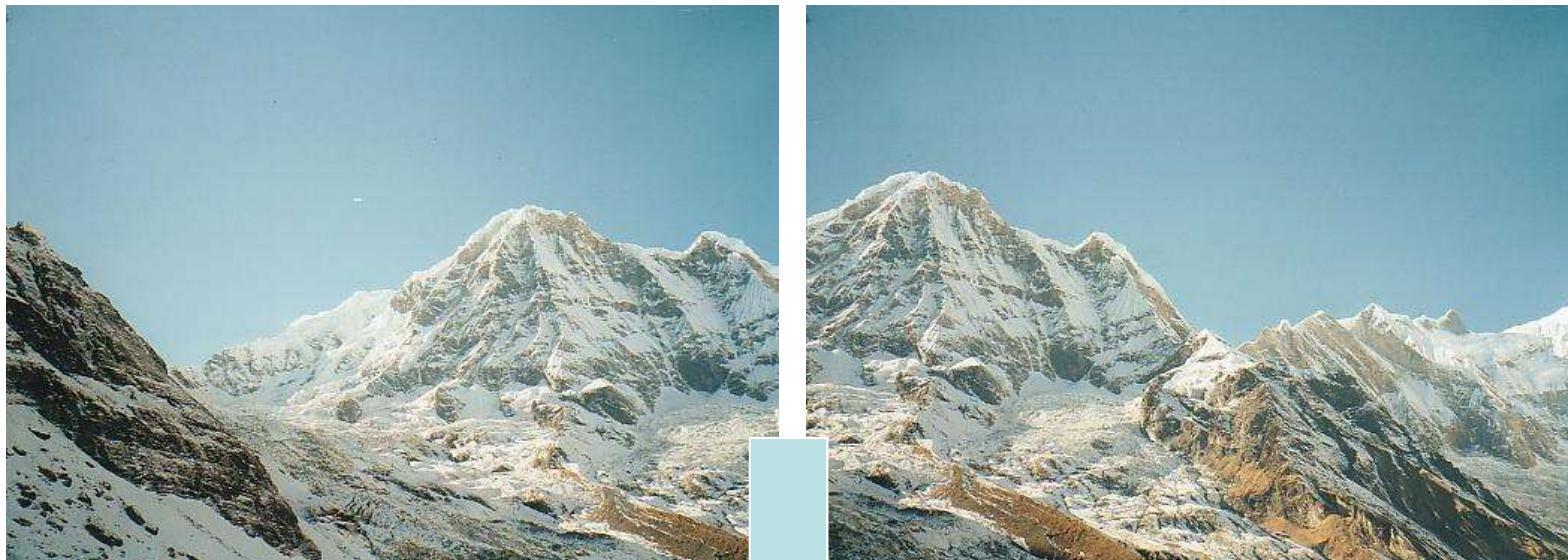
• • •
> 5000
images



8.3 Automatic image stitching

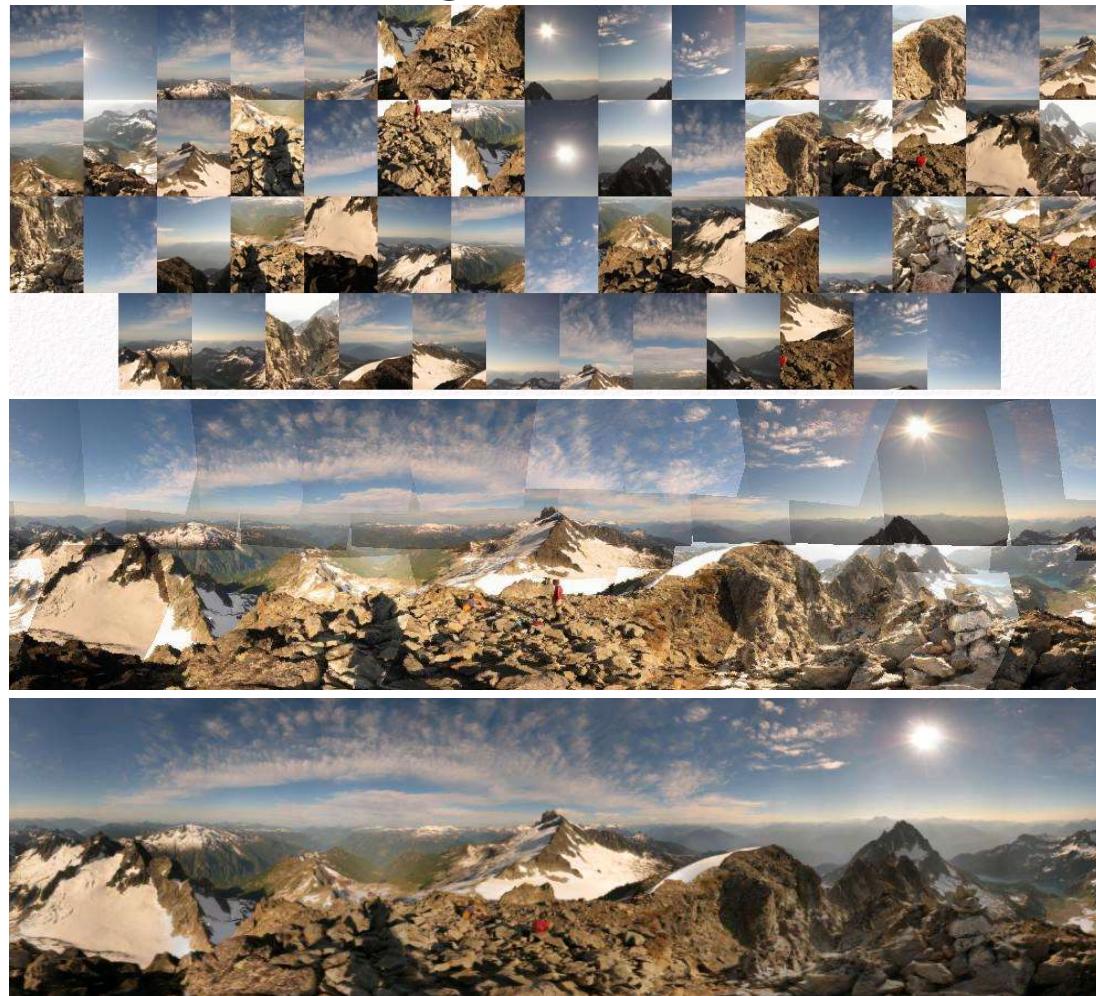


8.3 Automatic image stitching



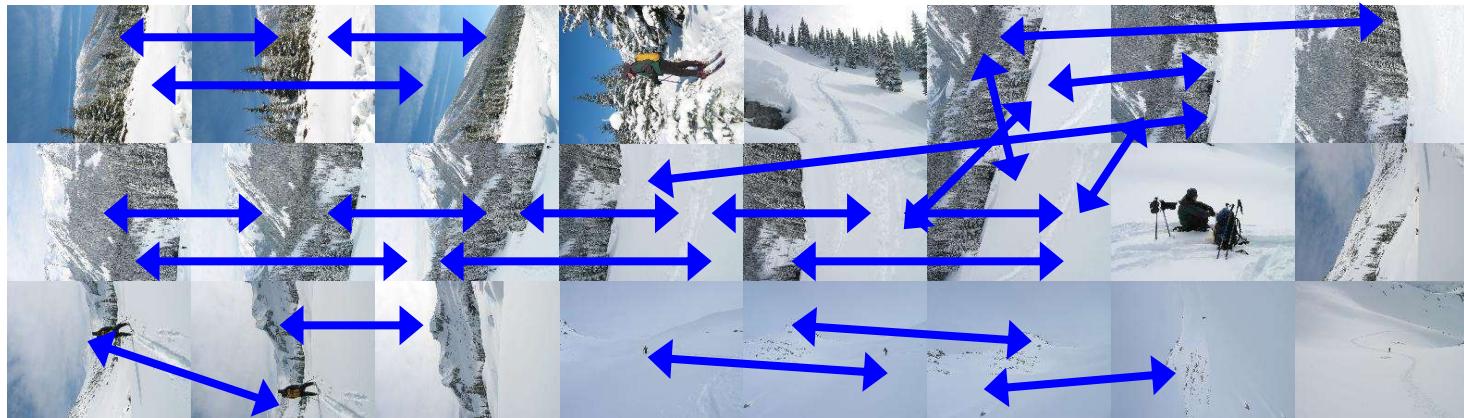
8.3 Image Stitch

- Automatic Mosaicing

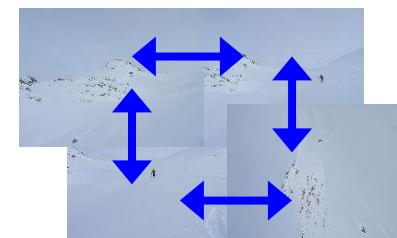
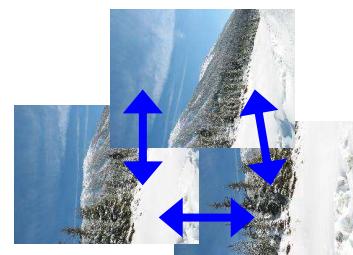
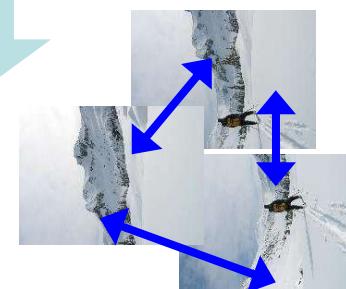
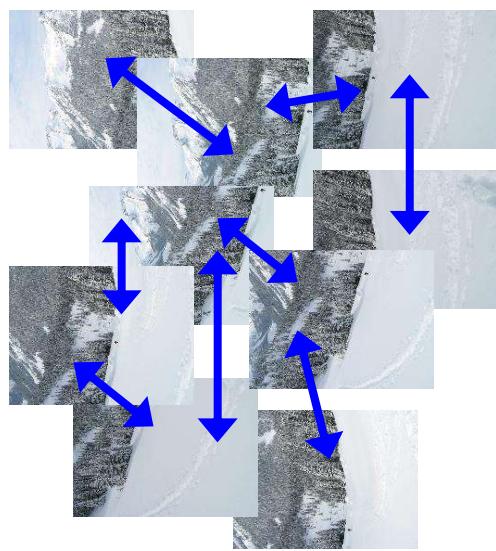
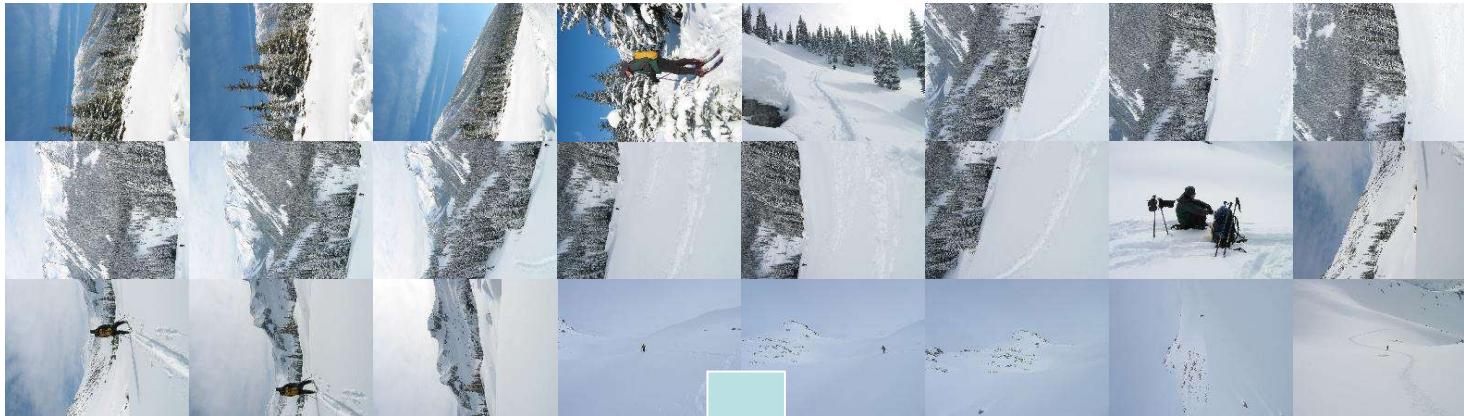


- <http://www.cs.ubc.ca/~mbrown/autostitch/autostitch.html>

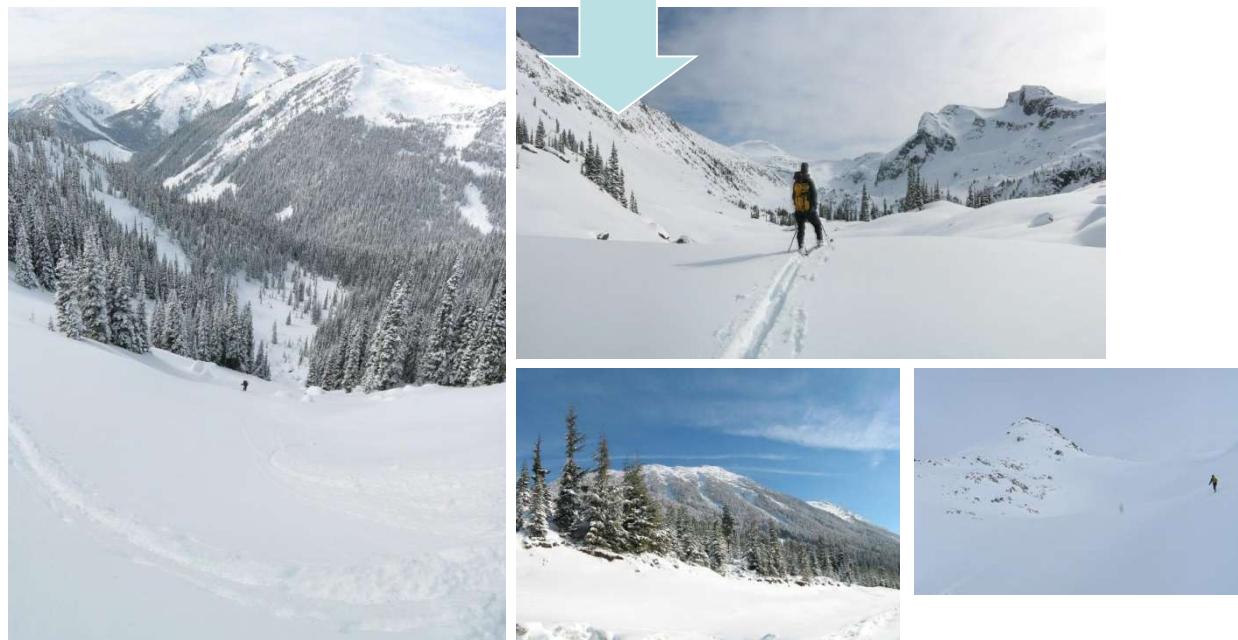
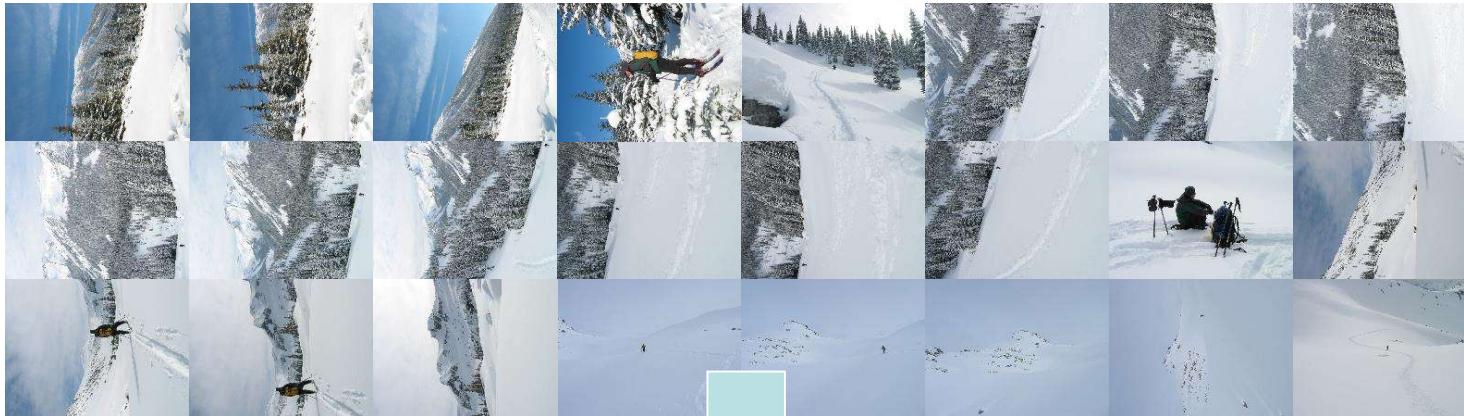
8.3 Automatic image stitching



8.3 Automatic image stitching



8.3 Automatic image stitching



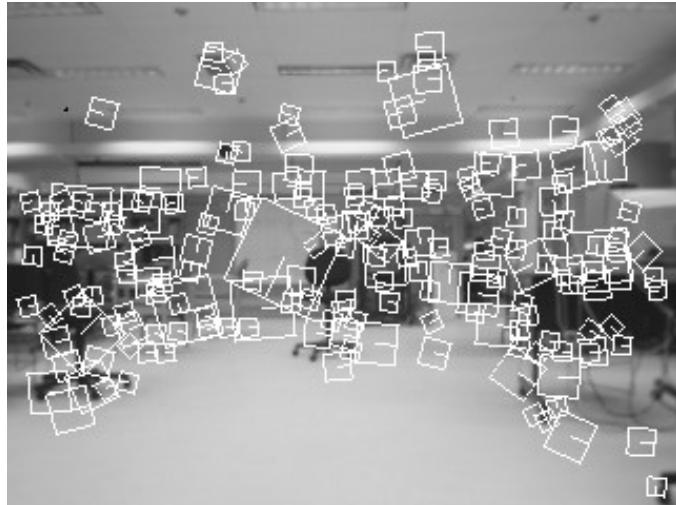
8.4 Robotics - Sony Aibo

SIFT is used for

- Recognizing charging station
- Communicating with visual cards
- Teaching object recognition
- soccer

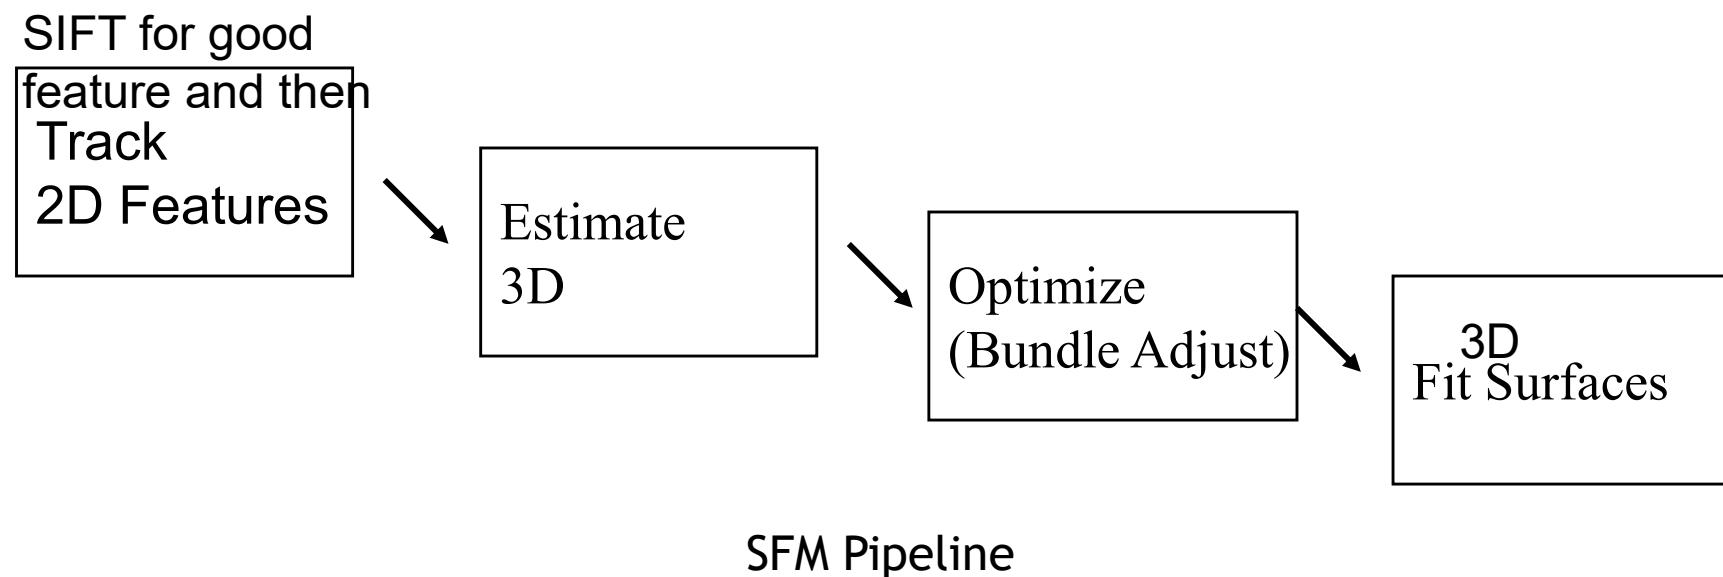


8.4 Robot location

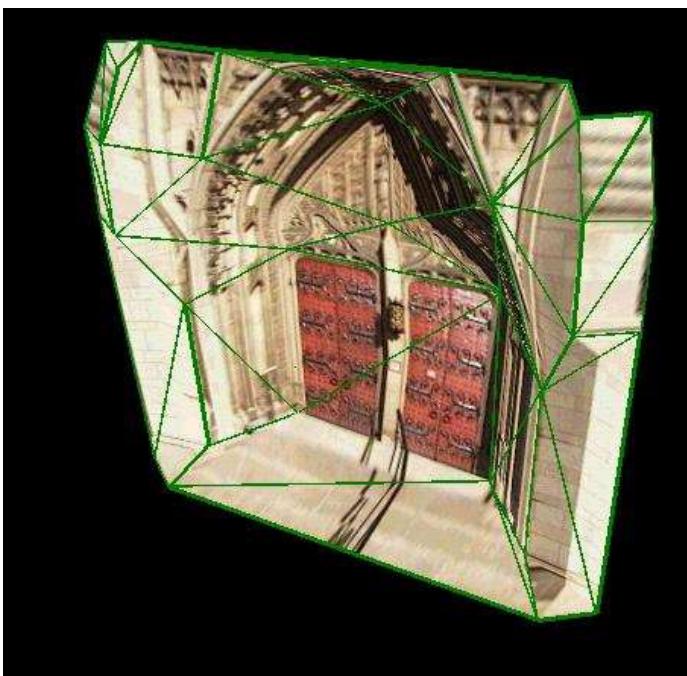
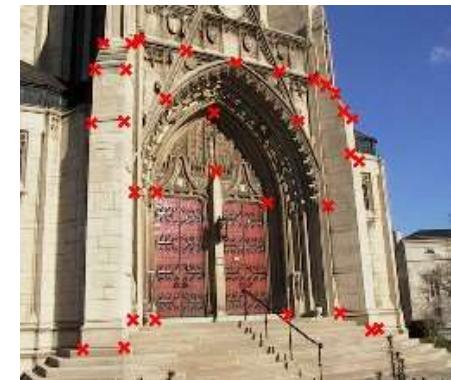
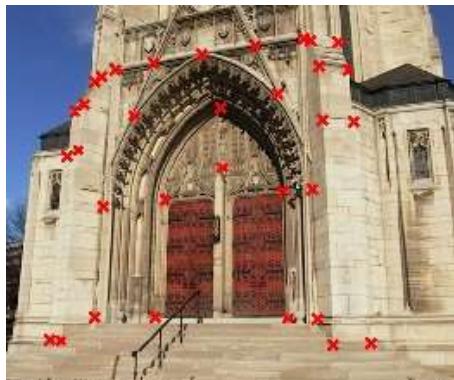


8.5 Structure from Motion (1/2)

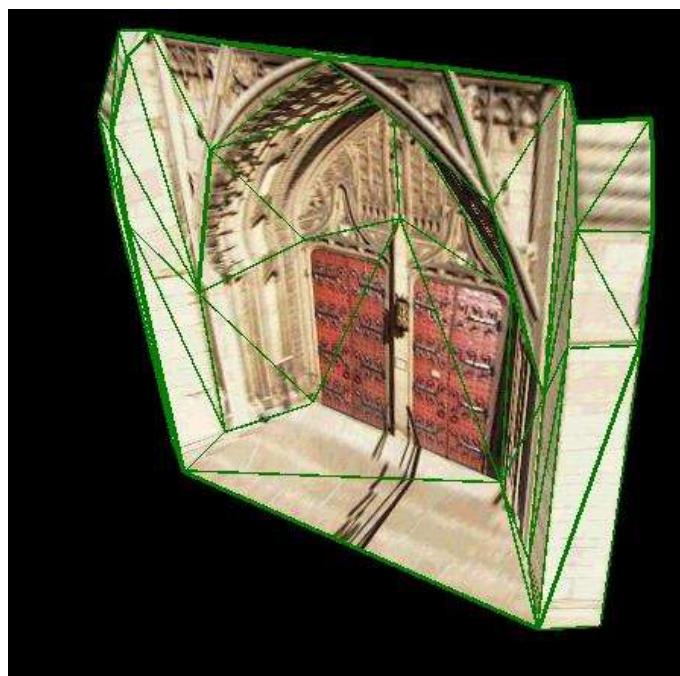
- The SFM Problem
 - Reconstruct scene geometry and camera motion from two or more images



8.5 Structure from Motion (2/2)



Poor mesh



Good mesh

8.6 Augmented reality



Reference

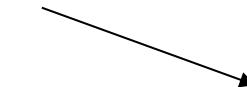
1. C. Harris and M. Stephens, “[A Combined Corner and Edge Detector](#),” 4th Alvey Vision Conference, pp147-151, 1988.
2. D.G. Lowe, “[Distinctive Image Features from Scale-Invariant Keypoints](#),” IJCV, 60(2), , pp91-110, 2004.

Append A. Gradient & Hessian of D

$$D(\mathbf{x}) = D + \frac{\partial D}{\partial \mathbf{x}}^T \mathbf{x} + \frac{1}{2} \mathbf{x}^T \frac{\partial^2 D}{\partial \mathbf{x}^2} \mathbf{x}$$



$$\begin{pmatrix} \frac{\partial f}{\partial x} \\ \frac{\partial f}{\partial y} \\ \frac{\partial f}{\partial \sigma} \end{pmatrix}$$



$$\begin{pmatrix} \frac{\partial^2 D}{\partial x^2} & \frac{\partial^2 D}{\partial x \partial y} & \frac{\partial^2 D}{\partial x \partial \sigma} \\ \frac{\partial^2 D}{\partial x \partial y} & \frac{\partial^2 D}{\partial y^2} & \frac{\partial^2 D}{\partial y \partial \sigma} \\ \frac{\partial^2 D}{\partial x \partial \sigma} & \frac{\partial^2 D}{\partial y \partial \sigma} & \frac{\partial^2 D}{\partial \sigma^2} \end{pmatrix}$$

$$\frac{\partial D}{\partial x} = \frac{D_{x+1,y,\sigma} - D_{x-1,y,\sigma}}{2}$$

$$\frac{\partial^2 D}{\partial x^2} = D_{x+1,y,\sigma} - 2D_{x,y,\sigma} + D_{x-1,y,\sigma}$$

$$\frac{\partial D}{\partial y} = \frac{D_{x,y+1,\sigma} - D_{x,y-1,\sigma}}{2}$$

$$\frac{\partial^2 D}{\partial y^2} = D_{x,y+1,\sigma} - 2D_{x,y,\sigma} + D_{x,y-1,\sigma}$$

$$\frac{\partial D}{\partial \sigma} = \frac{D_{x,y,\sigma+1} - D_{x,y,\sigma-1}}{2}$$

$$\frac{\partial^2 D}{\partial \sigma^2} = D_{x,y,\sigma+1} - 2D_{x,y,\sigma} + D_{x,y,\sigma-1}$$

$$\frac{\partial D}{\partial x \partial y} = \frac{D_{x+1,y+1,\sigma} - D_{x+1,y-1,\sigma} - D_{x-1,y+1,\sigma} + D_{x-1,y-1,\sigma}}{4}$$

$$\frac{\partial D}{\partial y \partial \sigma} = \frac{D_{x,y+1,\sigma+1} - D_{x,y+1,\sigma-1} - D_{x,y-1,\sigma+1} + D_{x,y-1,\sigma-1}}{4}$$

$$\frac{\partial D}{\partial x \partial \sigma} = \frac{D_{x+1,y,\sigma+1} - D_{x+1,y,\sigma-1} - D_{x-1,y,\sigma+1} + D_{x-1,y,\sigma-1}}{4}$$

Append B. Derivation of matrix form (1/2)

$$D(\mathbf{x}) = D + \boxed{\frac{\partial D^T}{\partial \mathbf{x}} \mathbf{x}} + \frac{1}{2} \mathbf{x}^T \frac{\partial^2 D}{\partial \mathbf{x}^2} \mathbf{x}$$

$$h(\mathbf{x}) = \boxed{\mathbf{g}^T \mathbf{x}} = \begin{pmatrix} g_1 & \cdots & g_n \end{pmatrix} \begin{pmatrix} x_1 \\ \vdots \\ x_n \end{pmatrix} = \sum_{i=1}^n g_i x_i$$

$$\frac{\partial h}{\partial \mathbf{x}} = \begin{pmatrix} \frac{\partial h}{\partial x_1} \\ \vdots \\ \frac{\partial h}{\partial x_n} \end{pmatrix} = \begin{pmatrix} g_1 \\ \vdots \\ g_n \end{pmatrix} = \mathbf{g}$$

Append B. Derivation of matrix form (2/2)

$$D(\mathbf{x}) = D + \frac{\partial D^T}{\partial \mathbf{x}} \mathbf{x} + \frac{1}{2} \boxed{\mathbf{x}^T \frac{\partial^2 D}{\partial \mathbf{x}^2} \mathbf{x}}$$

$$h(\mathbf{x}) = \boxed{\mathbf{x}^T \mathbf{A} \mathbf{x}} = (x_1 \quad \dots \quad x_n)^T \begin{pmatrix} a_{11} & \cdots & a_{1n} \\ \vdots & \ddots & \vdots \\ a_{n1} & \cdots & a_{nn} \end{pmatrix} \begin{pmatrix} x_1 \\ \vdots \\ x_n \end{pmatrix}$$

$$= \sum_{i=1}^n \sum_{j=1}^n a_{ij} x_i x_j$$

$$\frac{\partial h}{\partial \mathbf{x}} = \begin{pmatrix} \frac{\partial h}{\partial x_1} \\ \vdots \\ \frac{\partial h}{\partial x_n} \end{pmatrix} = \begin{pmatrix} \sum_{i=1}^n a_{i1} x_i + \sum_{j=1}^n a_{1j} x_j \\ \vdots \\ \sum_{i=1}^n a_{in} x_i + \sum_{j=1}^n a_{nj} x_j \end{pmatrix} = \mathbf{A}^T \mathbf{x} + \mathbf{A} \mathbf{x}$$

$$= (\mathbf{A}^T + \mathbf{A}) \mathbf{x}$$